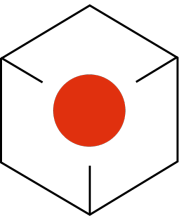


# Dynamische Attribute zur Speicherung von flexiblen Daten

Dr. Andrea Kennel  
InfoPunkt Kennel GmbH  
Dübendorf-Schweiz  
November 2023



# Dr. Andrea Kennel

**SYM<sup>L2</sup>**



Consultant

Dozentin für Datenbanken

Coach für Project Management

Fachhochschule Nordwestschweiz

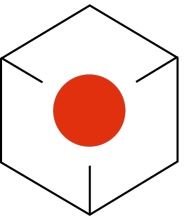
Brugg/Windisch, Schweiz



[andrea.kennel@fhnw.ch](mailto:andrea.kennel@fhnw.ch)

[andrea@infokennel.ch](mailto:andrea@infokennel.ch)

[www.infokennel.ch](http://www.infokennel.ch)

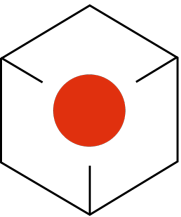


# Ausgangslage

Ein neues Start-Up will einen **Web-Shop** für **Netzwerkgeräte** aufbauen. Die Daten zu diesen Geräten sollen in einer Datenbank verwaltet werden.

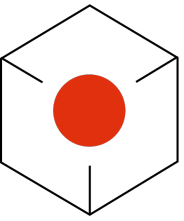
Der Web-Shop wird mit Java-Skript programmiert und die Daten zu den Netzwerkgeräten werden als **JSON** geliefert.

User Job ist das Erstellen der Datenbank.



```
{  "type": "router",
  "name": "Hans B4",
  "description": "4 port router",
  "manufacturer": "XYZ",
  "price": 800,
  "port_group": [
    { "amount": 4,
      "type": "RJ45",
      "speeds": "100/1000"
    },
    { "amount": 2,
      "type": "SFP",
      "speeds": "1000/10000"
    }
  ],
  "routing": {
    "protocols": "static OSPF",
    "table_size": 5
  }
}
```

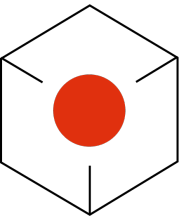




# Wie sehen die Geräte aus?

```
{  "type": "router",
    "name": "Hans B4",
    "description": "4 port router",
    "manufacturer": "XYZ",
    "price": 800,
    "port_group": [
      { "amount": 4,
        "type": "RJ45",
        "speeds": "100/1000"
      },
```

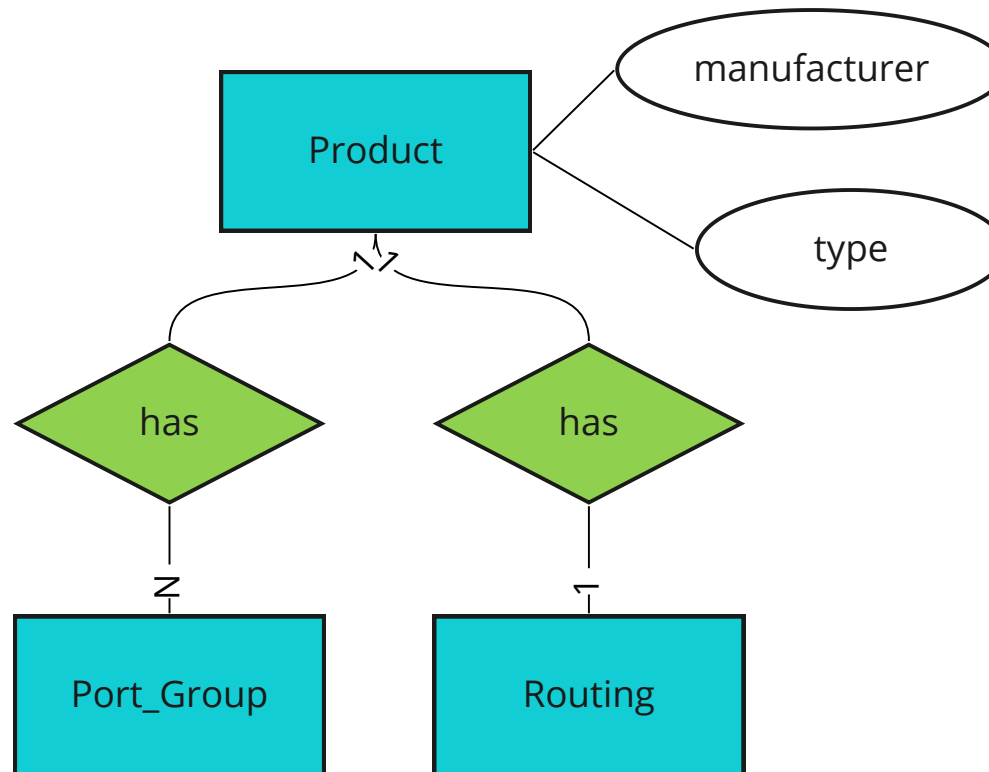
```
    { "amount": 2,
      "type": "SFP",
      "speeds": "1000/10000"
    },
    "routing": {
      "protocols": "static OSPF",
      "table_size": 5
    }
}
```

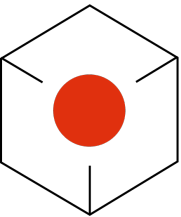


# Kannst Du mal schnell

*Für einen Web-shop, der Netzwerkgeräte verkauft, müssen die Daten zu den Netzwerkgräten in einer Datenbank gespeichert werden.*

*Die Daten liegen als JSON vor, das Modell ist relativ einfach:*





# Datenmodell umsetzen und Daten erfassen

```
CREATE TABLE product (
  prod_id    NUMBER(10),
  type       VARCHAR2(50),
  name       VARCHAR2(50),
  description VARCHAR2(200),
  manufacturer VARCHAR2(50),
  price      NUMBER (8,2)
);

CREATE TABLE port_group (
```

Query Result x Script Output x  
Task completed in 1.972 seconds

Table PRODUCT created.

Table PORT\_GROUP created.

Table ROUTING created.

```
INSERT INTO product (prod_id, type, name, description, manufacturer
VALUES (1, 'network switch', 'Fritz A16', 'unmanaged 16 port network
(2, 'network switch', 'Fritz B24', 'unmanaged 24 port 10gbit
(3, 'network switch', 'Fritz C12', 'managed 12 port network
(4, 'router', 'Hans A8', '8 port router', 'ABC', 1000),
(5, 'router', 'Hans B4', '4 port router', 'XYZ', 800);

INSERT INTO port_group (pogr_id, prod_id, amount, type, speeds)...

INSERT INTO routing (rout_id, prod_id, protocols, table_size)
```

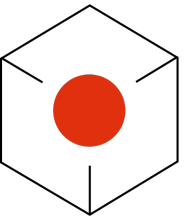
Query Result x Script Output x  
Task completed in 0.597 seconds

5 rows inserted.

8 rows inserted.

2 rows inserted.



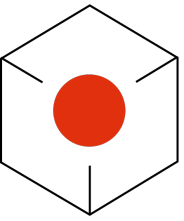


# Wir haben da ein neues Gerät

```
{ "type": "network switch/layer3 switch",  
  "name": "Fritz C16",  
  "description": "16 port PoE layer 3 network switch",  
  "manufacturer": "ABC",  
  "price": 500,  
  "port_group": [  
    { "amount": 16,  
      "type": "RJ45",  
      "speeds": "10/100/1000",  
      "poe": {"modes": ["active", "passive"],  
              "volt": [24, 48]}  
    }  
  ],  
  ...
```

```
    "feature": [  
      { "name": "VLAN",  
        "amount": 4094},  
      { "name": "QoS",  
        "amount": 8},  
      { "name": "network access control",  
        "type": "MAC based authentication",  
        "vlan_support": true},  
      { "name": "routing",  
        "protocols": "static, RIP, OSPF, BGP",  
        "table_size": 10}  
    ]  
  }
```





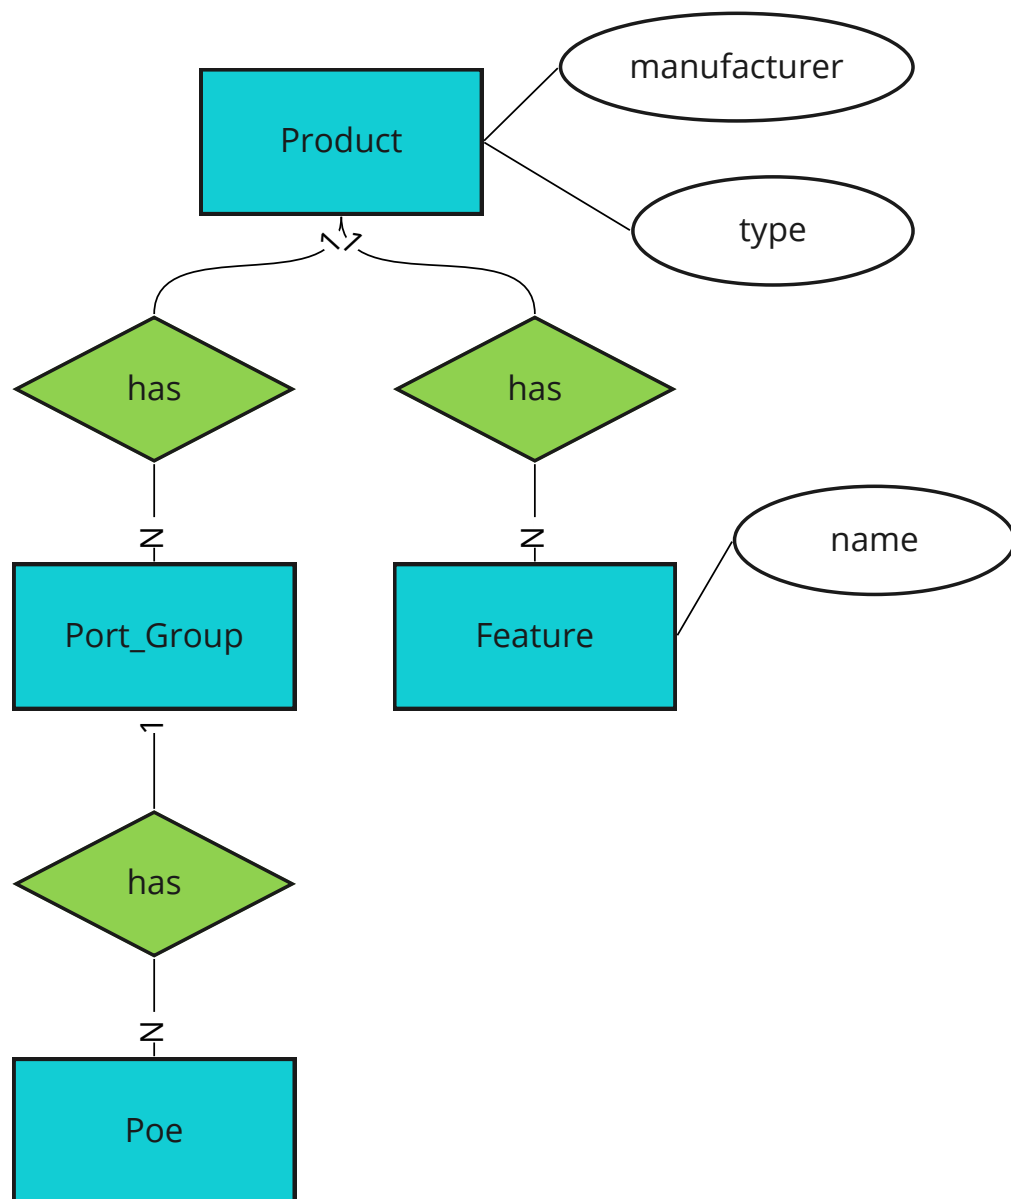
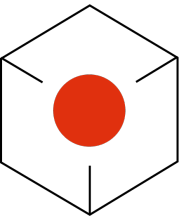
# Dann müssen wir die Datenstruktur anpassen

Detail-Tabelle zu Port-Group und

Feature als Generalisierung mit mehreren Spezialisierungen

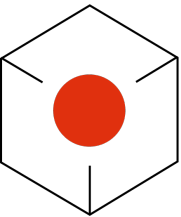
Wir wissen nicht, was noch kommt und fassen alle Attribute in der Generalisierung zusammen.





# Neue Datenstruktur





# Neue Datenstruktur

Worksheet | Query Builder

```
CREATE TABLE poe (  
  poe_id      NUMBER(10),  
  pogr_id     NUMBER(10),  
  modes       VARCHAR2(20),  
  volt        NUMBER (6,2)  
);  
  
ALTER TABLE poe ADD CONSTRAINT poe_pk PRIMARY KEY (poe_id);  
ALTER TABLE poe ADD CONSTRAINT poe_pogr_fk FOREIGN KEY (pogr_id)  
REFERENCES port_group (pogr_id);
```

Query Result | Script Output

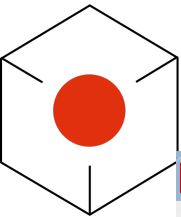
Task completed in 0.555 seconds

Table P0E created.

Table P0E altered.

Table P0E altered.





# Datenstruktur

Worksheet | Query Builder

```
RENAME routing TO feature;  
ALTER TABLE feature RENAME COLUMN rout_id TO feat_id;  
ALTER TABLE feature ADD name VARCHAR2(200);  
ALTER TABLE feature ADD amount NUMBER(10);  
ALTER TABLE feature ADD type VARCHAR2(200);  
ALTER TABLE feature ADD vlan_support VARCHAR2(200);  
-- all existing get name 'routing'  
UPDATE feature  
SET name = 'routing';
```

Query Result | Script Output

Task completed in 1.832 seconds

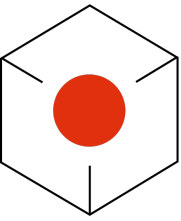
Table FEATURE altered.

Table FEATURE altered.

2 rows updated.







# Was, wenn es mehr Attribute gibt?

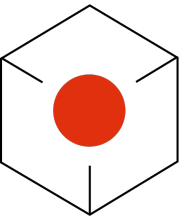
Wir mussten bei FEATURES weitere Attribute zufügen.

Was passiert, wenn da noch weitere Attribute kommen?

Was wenn viele features dann nur wenige Attribute gefüllt haben?

Gibt es da eine flexiblere Lösung?

```
"feature": [  
  { "name": "VLAN",  
    "amount": 4094},  
  { "name": "QoS",  
    "amount": 8},  
  { "name": "network access control",  
    "type": "MAC based authentication",  
    "vlan_support": true},  
  { "name": "routing",  
    "protocols": "static, RIP, OSPF, BGP",  
    "table_size": 10}  
]  
}
```



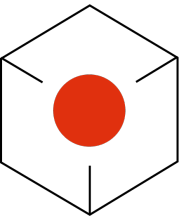
# Was, wenn es mehr Attribute gibt?

Die Antwort heisst EAV

Das Entity-Attribute-Value-Modell (EAV) ist eine Datenmodellierungstechnik, die in Datenbanken verwendet wird, um Daten auf flexible und skalierbare Weise zu speichern und abzurufen.

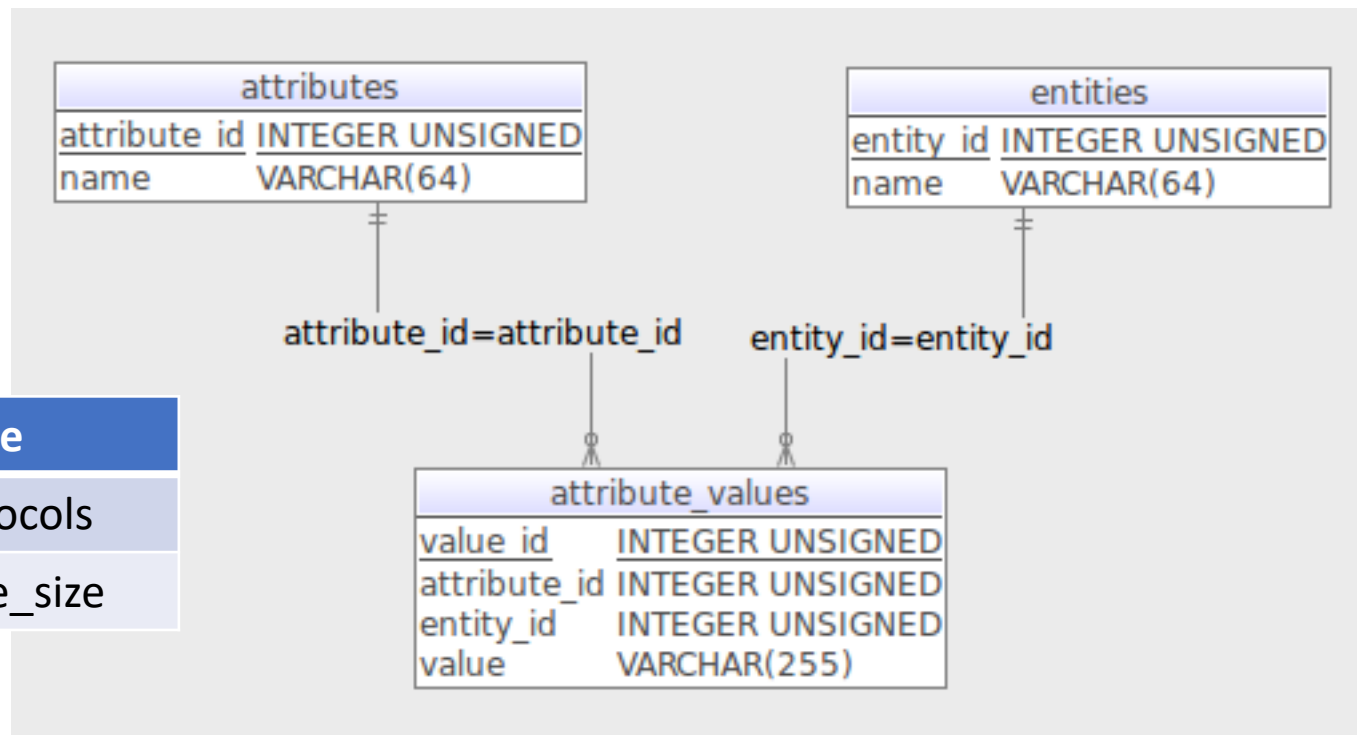
Details zum Modell: <https://inviqa.com/blog/understanding-eav-data-model-and-when-use-it>

**EAV**



# EAV

Attribute_id	name
84	protocols
85	table_size

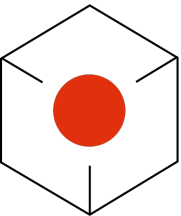


```
INSERT INTO feature (feat_id, prod_id, name, protocols, table_size)
VALUES (24, 21, 'routing', 'static, RIP, OSPF, BGP', 10);
```

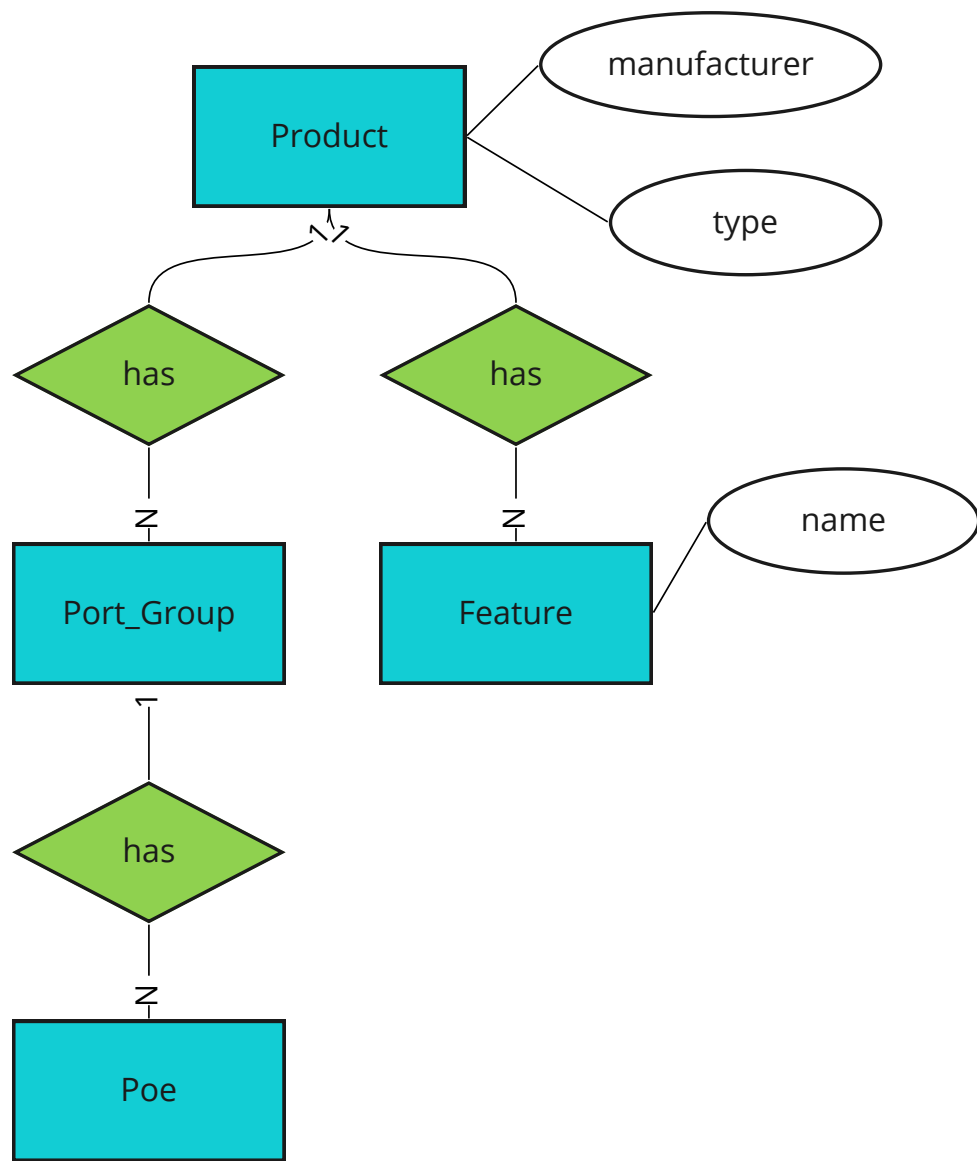
Value_id	Attribute_id	Entity_id	values
1	84	24	static, RIP, OSPF, BGP
2	85	24	10

# EAV

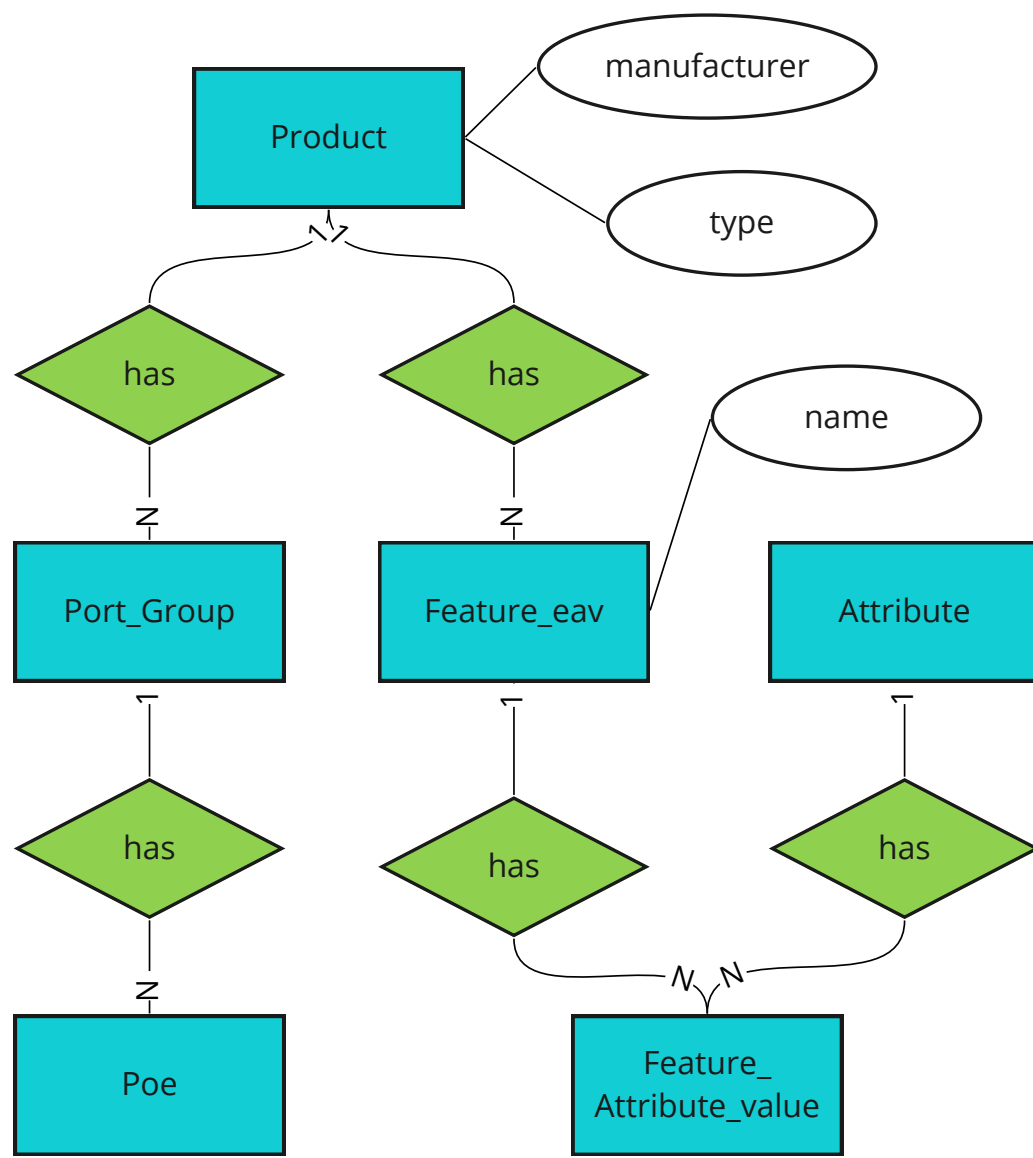


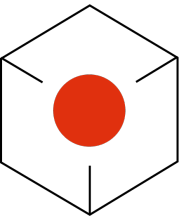


EAV



EAV





# EAV

```
Welcome Page | Docker_loacal_ANDREA | edbsHS2023_AN

Worksheet | Query Builder

CREATE TABLE feature_eav (
  feat_id      NUMBER(10),
  prod_id      NUMBER(10),
  name         VARCHAR2(200)
);

CREATE TABLE attribute (
  attr_id      NUMBER(10),
  name         VARCHAR2(30)
);

CREATE TABLE feature_attribute_value (
  valu_id      NUMBER(10),
  feat_id      NUMBER(10),
  attr_id      NUMBER(10),
  value        VARCHAR2(2000)
);
```

```
Welcome Page | Docker_loacal_ANDREA | edbsHS2023_ANDREA | Docker_loacal_A

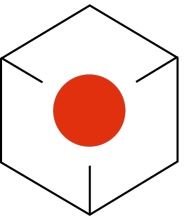
Worksheet | Query Builder

INSERT INTO feature_eav (feat_id, prod_id, name)
VALUES (24, 21, 'routing');

INSERT INTO attribute (attr_id, name)
VALUES (84, 'protocols'),
       (85, 'table_size');

INSERT INTO feature_attribute_value (valu_id, feat_id, attr_id, value)
VALUES (25, 24, 84, 'static, RIP, OSPF, BGP'),
       (26, 24, 85, '10');
```

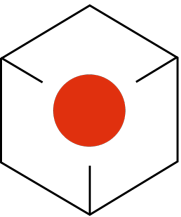
# EAV



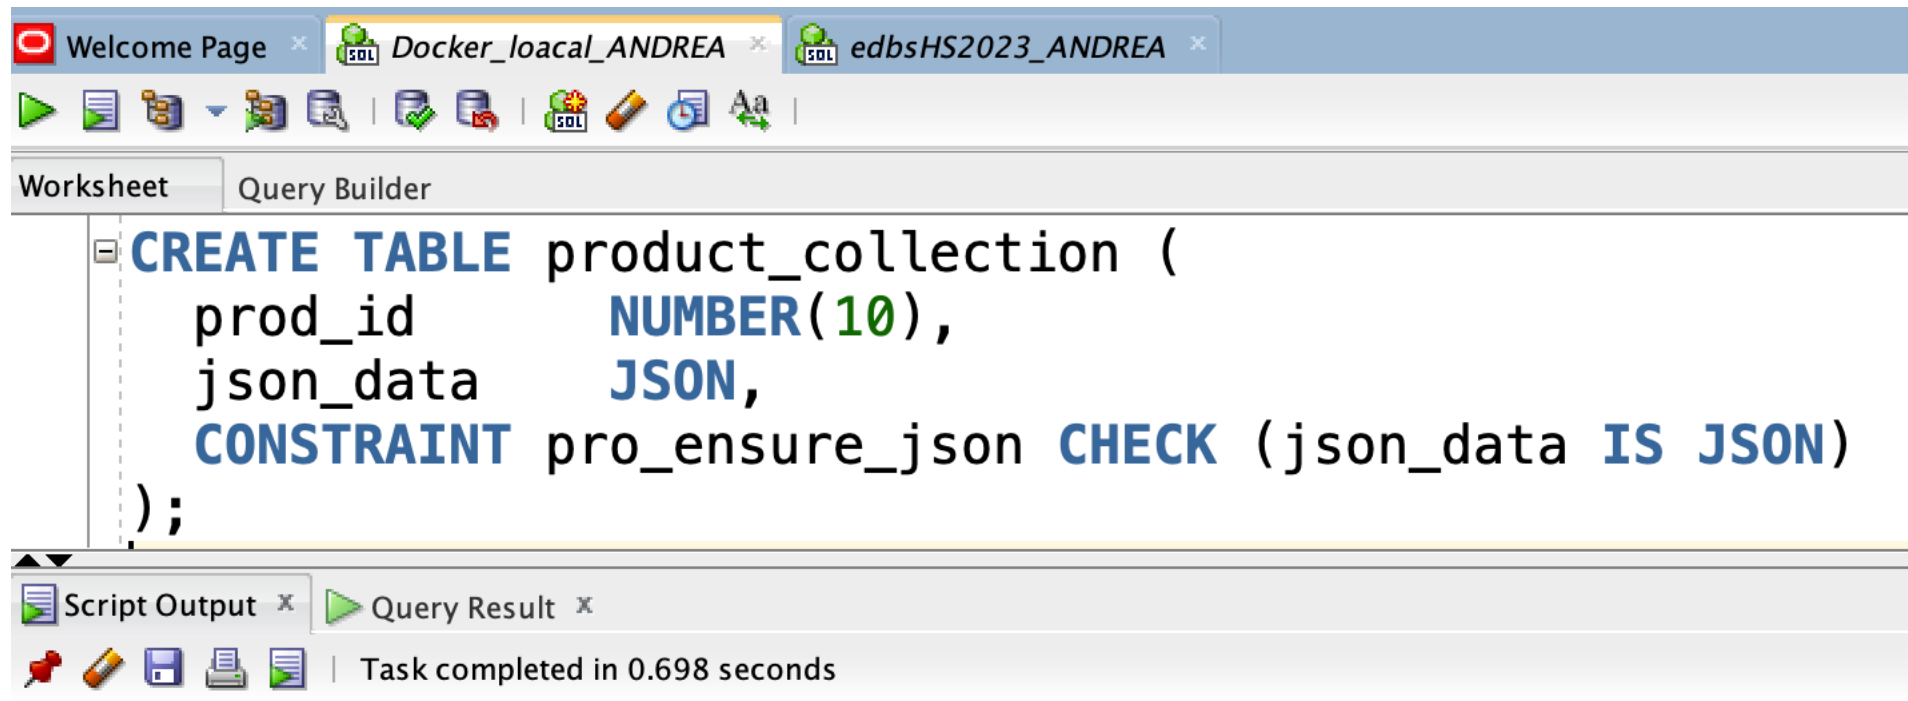
# Wäre da nicht eine Dokument DB besser?

- Ich verliere den Überblick
- Warum nicht einfach alle Daten als JSON speichern?





# Wäre da nicht eine Dokument DB besser?



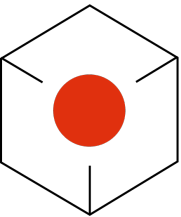
The screenshot shows a SQL IDE window with three tabs: 'Welcome Page', 'Docker\_loacal\_ANDREA', and 'edbsHS2023\_ANDREA'. The 'Query Builder' tab is active, displaying the following SQL code:

```
CREATE TABLE product_collection (  
  prod_id      NUMBER(10),  
  json_data    JSON,  
  CONSTRAINT pro_ensure_json CHECK (json_data IS JSON)  
);
```

Below the code editor, the 'Script Output' and 'Query Result' tabs are visible. The status bar at the bottom indicates 'Task completed in 0.698 seconds'.



Table PRODUCT\_COLLECTION created.



# Wäre da nicht eine Dokument DB besser?

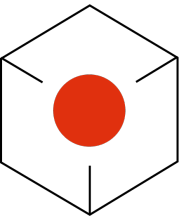
Two screenshots of a SQL IDE interface comparing SQL and JSON data storage.

**Left Screenshot:** The 'Worksheet' tab is active, showing an SQL `INSERT` statement for a `product_collection` table. The JSON data is stored as a string within the `json_data` column.

```
INSERT INTO product_collection (prod_id, json_data)
VALUES (1, json('{
  "type": "network switch",
  "name": "Fritz A16",
  "description": "unmanaged 16 port netwo
  "manufacturer": "ABC",
  "price": 100,
  "port_group": [
    {
      "amount": 16,
      "type": "RJ45",
      "speeds": "10/100/1000"
    }
  ]
}'));
```

**Right Screenshot:** The 'Query Builder' tab is active, showing the same data structure as a JSON object, which is more readable and structured than the SQL string.

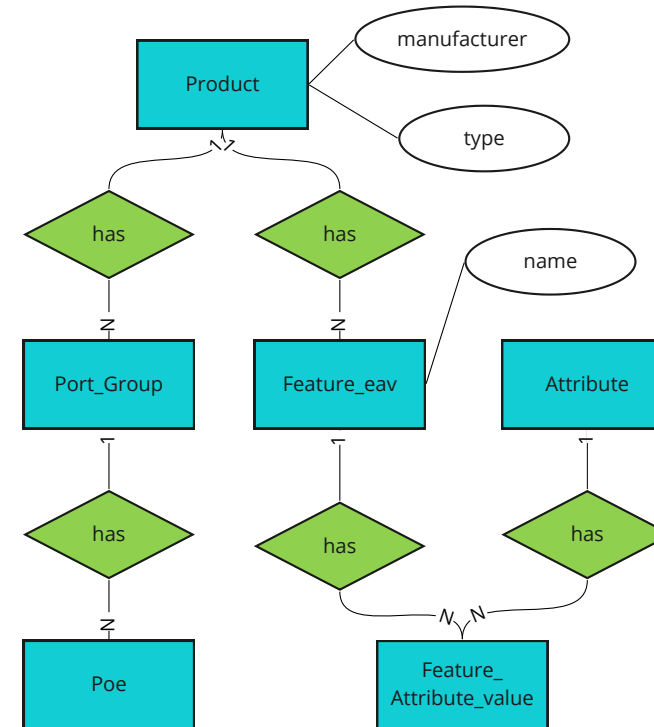
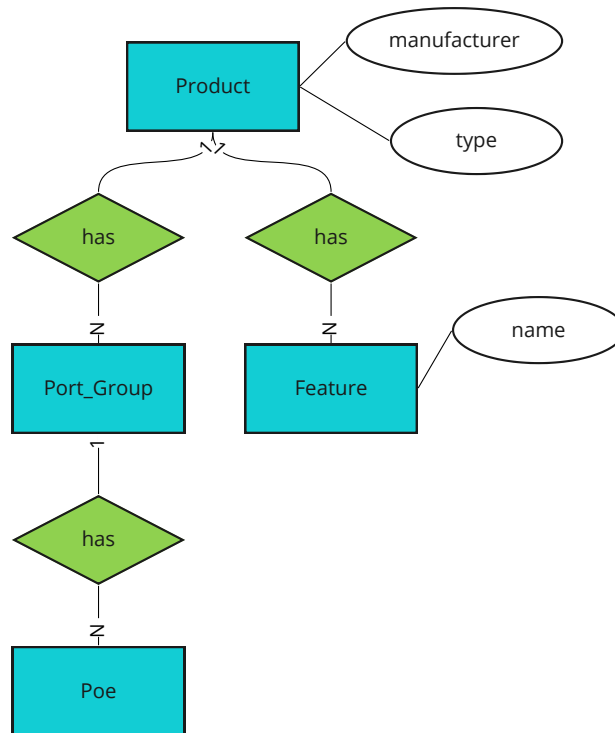
```
INSERT INTO product_collection (prod_id, json_data)
VALUES (21, json({
  "type": "network switch/layer3 switch",
  "name": "Fritz C16",
  "description": "16 port PoE layer 3 network switch",
  "manufacturer": "ABC",
  "price": 500,
  "port_group": [
    {
      "amount": 16,
      "type": "RJ45",
      "speeds": "10/100/1000",
      "poe": {
        "modes": [
          "active",
          "passive"
        ],
        "volt": [
          24,
          48
        ]
      }
    }
  ],
  "feature": [
    {
      "name": "VLAN",
      "amount": 4094
    },
    {
      "name": "QoS",
      "amount": 8
    },
    {
      "name": "network access control",
      "type": "MAC based authentication",
      "vlan_support": true
    },
    {
      "name": "routing",
      "protocols": "static, RIP, OSPF, BGP",
      "table_size": 10
    }
  ]
}'));
```

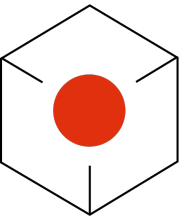


# Wäre da nicht eine Dokument DB besser?

Wie sieht das mit Abfragen aus?

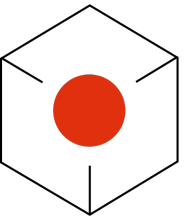
Beispiel: Haben wir Geräte mit einer Port-Group vom Typ SFP?





# Haben wir Geräte mit einer Port-Group vom Typ SFP? relational

```
SELECT pro.prod_id, pro.type product_type, pro.name, pog.amount, pog.type
FROM product pro INNER JOIN
    port_group pog ON (pro.prod_id = pog.prod_id)
WHERE pog.type = 'SFP';
```



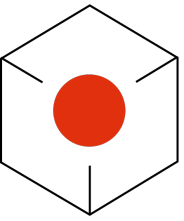
# Haben wir Geräte mit einer Port-Group vom Typ SFP?

## JSON

```
SELECT pc.prod_id, pc.json_data.type product_type,  
       pc.json_data.name,  
       pc.json_data.port_group[*].amount,  
       pc.json_data.port_group[*].type  
FROM product_collection pc  
WHERE json_exists(pc.json_data,  
                  '$.port_group?(@.type == $v1)'  
                  PASSING 'SFP' AS "v1");
```

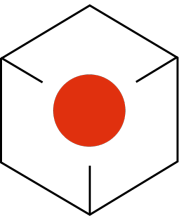
```
SELECT pc.prod_id, pc.json_data.type product_type,  
       pc.json_data.name,  
       jt.*  
FROM product_collection pc,  
     JSON_TABLE(  
       pc.json_data  
       COLUMNS (  
         NESTED port_group[*]  
         COLUMNS (  
           amount NUMBER(3) PATH '$.amount',  
           type VARCHAR2(50) PATH '$.type'  
         )  
       )  
     ) jt  
WHERE jt.type = 'SFP';
```





Haben wir Geräte mit Feature, das amount = 8 hat?  
relational

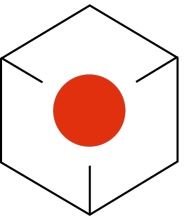
```
SELECT pro.prod_id, pro.type product_type, pro.name,  
       fe.name feature_name, fe.amount  
FROM product pro INNER JOIN  
       feature fe ON (pro.prod_id = fe.prod_id)  
WHERE fe.amount = 8;
```



# Haben wir Geräte mit Feature, das amount = 8 hat?

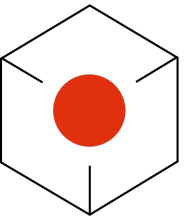
## JSON

```
SELECT pc.prod_id, pc.json_data.type product_type, pc.json_data.name,  
       jt.*  
FROM product_collection pc,  
     JSON_TABLE(  
       pc.json_data  
       COLUMNS (  
         NESTED feature[*]  
         COLUMNS (  
           feature_name VARCHAR2(50) PATH '$.name',  
           amount NUMBER(10) PATH '$.amount'  
         )  
       )  
     ) jt  
WHERE jt.amount = 8;
```

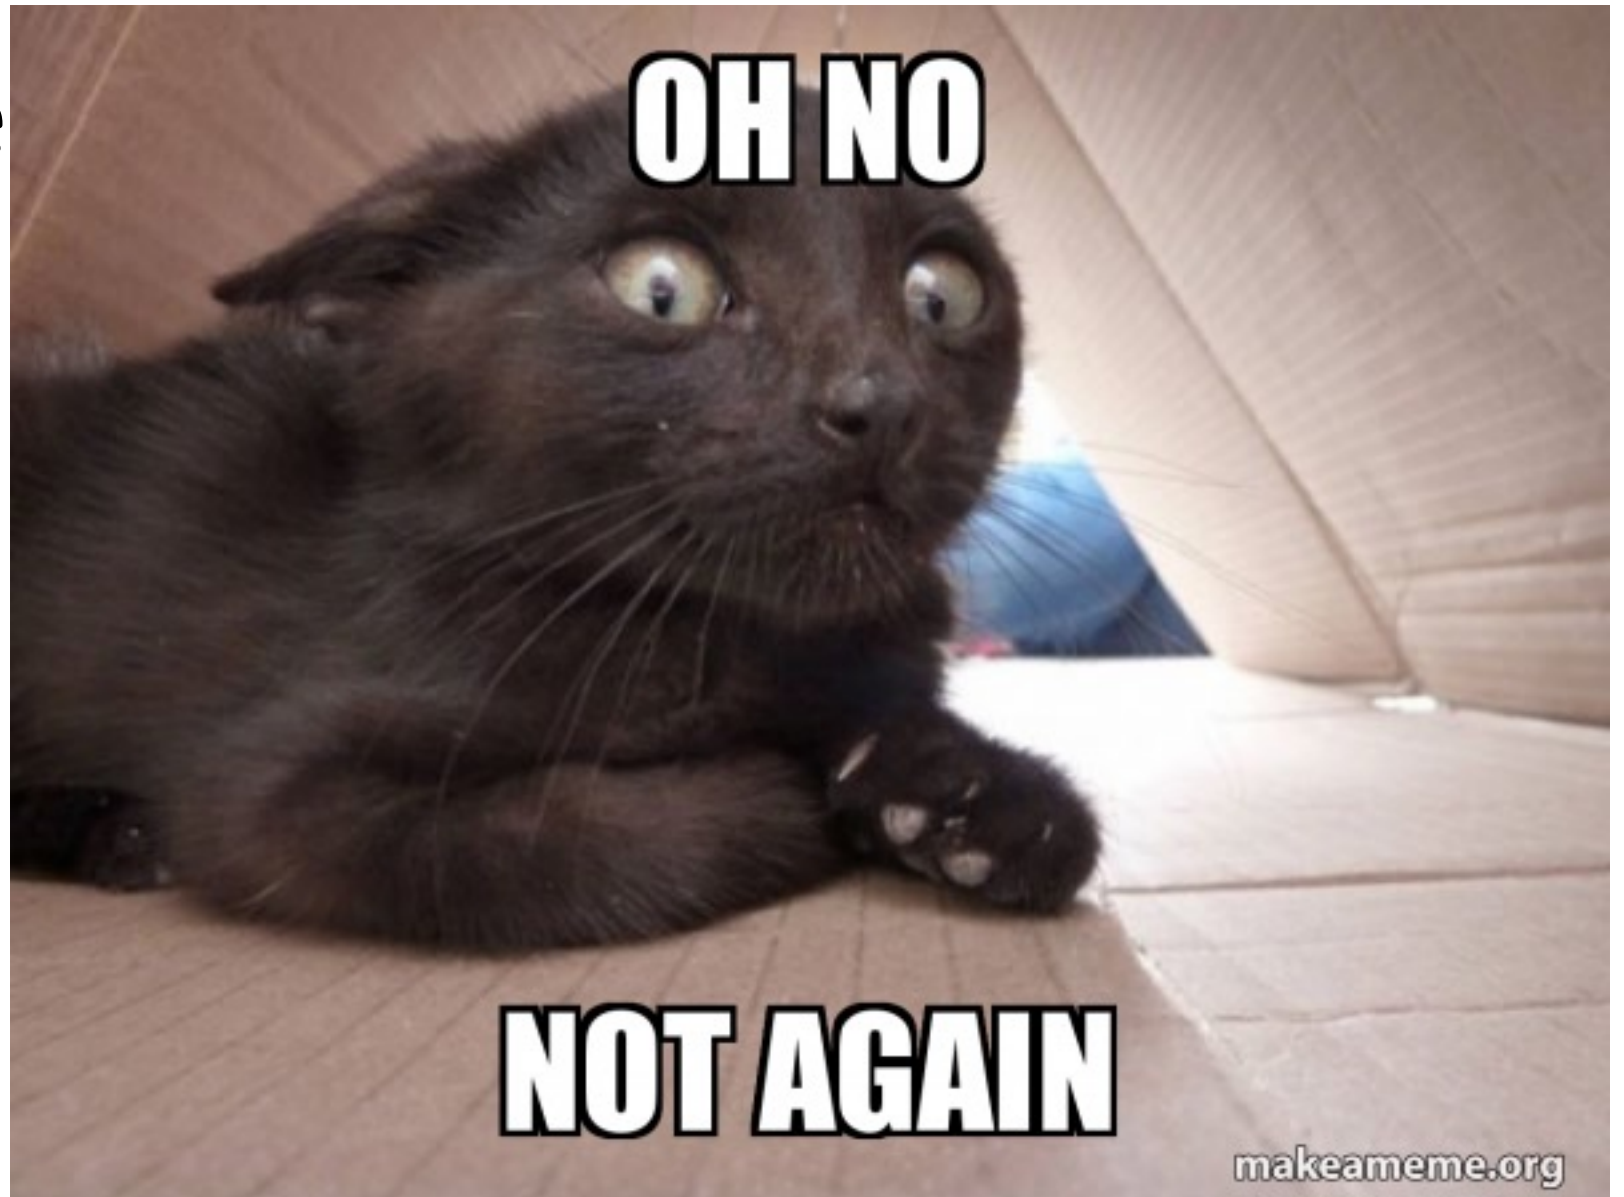


Haben wir Geräte mit Feature, das amount = 8 hat?  
EAV

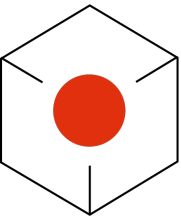
```
SELECT pro.prod_id, pro.type product_type, pro.name,  
       fe.name feature_name, feat.value amount  
FROM product pro INNER JOIN  
     feature_eav fe ON (pro.prod_id = fe.prod_id) INNER JOIN  
     feature_attribute_value feat ON (fe.feat_id = feat.feat_id) INNER JOIN  
     attribute at ON (feat.attr_id = at.attr_id)  
WHERE at.name = 'amount'  
AND feat.value = '8';
```



Wir haben wie



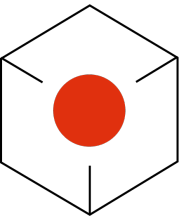
<https://makeameme.org/meme/oh-no-not-6bc4c7>



# Wir haben wieder ein neues Gerät

```
{ "type": "router/firewall",  
  "name": "Hans A48",  
  ....  
  "feature": [  
    { ... },  
    { ... },  
    { ... },  
    { "name": "firewall",  
      "type": "stateful, stateless",  
      "amount_of_rules": 1000,  
      "amount_of_connections": 100000,  
      "amount_of_nat_rules": 1000,  
      "feature": [  
        { "name": "URL filtering",  
          "types": "whitelist, blacklist",  
          "amount_of_rules": 1000 },  
        ....  
      ]  
    }  
  ]  
}
```

```
{ "name": "application filtering",  
  "types": "whitelist",  
  "amount_of_rules": 1000 },  
{ "name": "content filtering",  
  "types": "blacklist",  
  "amount_of_rules": 1000 },  
{ "name": "anti-virus",  
  "ssl_inspection": true }  
]  
{ "name": "VPN",  
  "type": "IPsec, SSL, L2TP, PPTP",  
  "amount_of_tunnels": 1000,  
  "amount_of_users": 1000,  
  "feature": [  
    { "name": "IPsec",  
      "type": "IKEv1, IKEv2",  
      "amount_of_tunnels": 1000,  
      "amount_of_users": 1000 },  
    { "name": "SSL",  
      ....  
    }  
  ]  
}
```



# Dann müssen wir die Datenstruktur anpassen

Weiter Typen von Features mit neuen Attributen, damit haben wir ja schon gerechnet

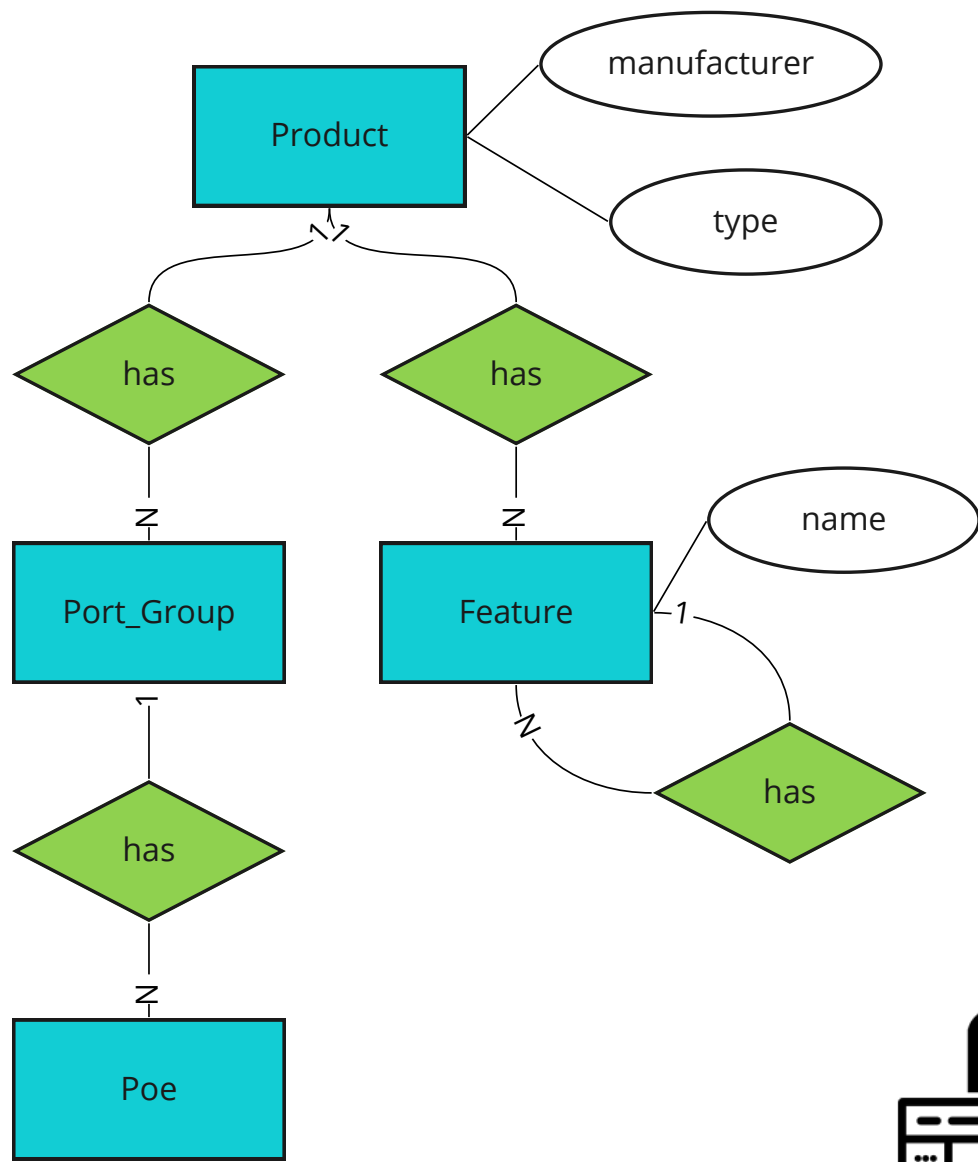
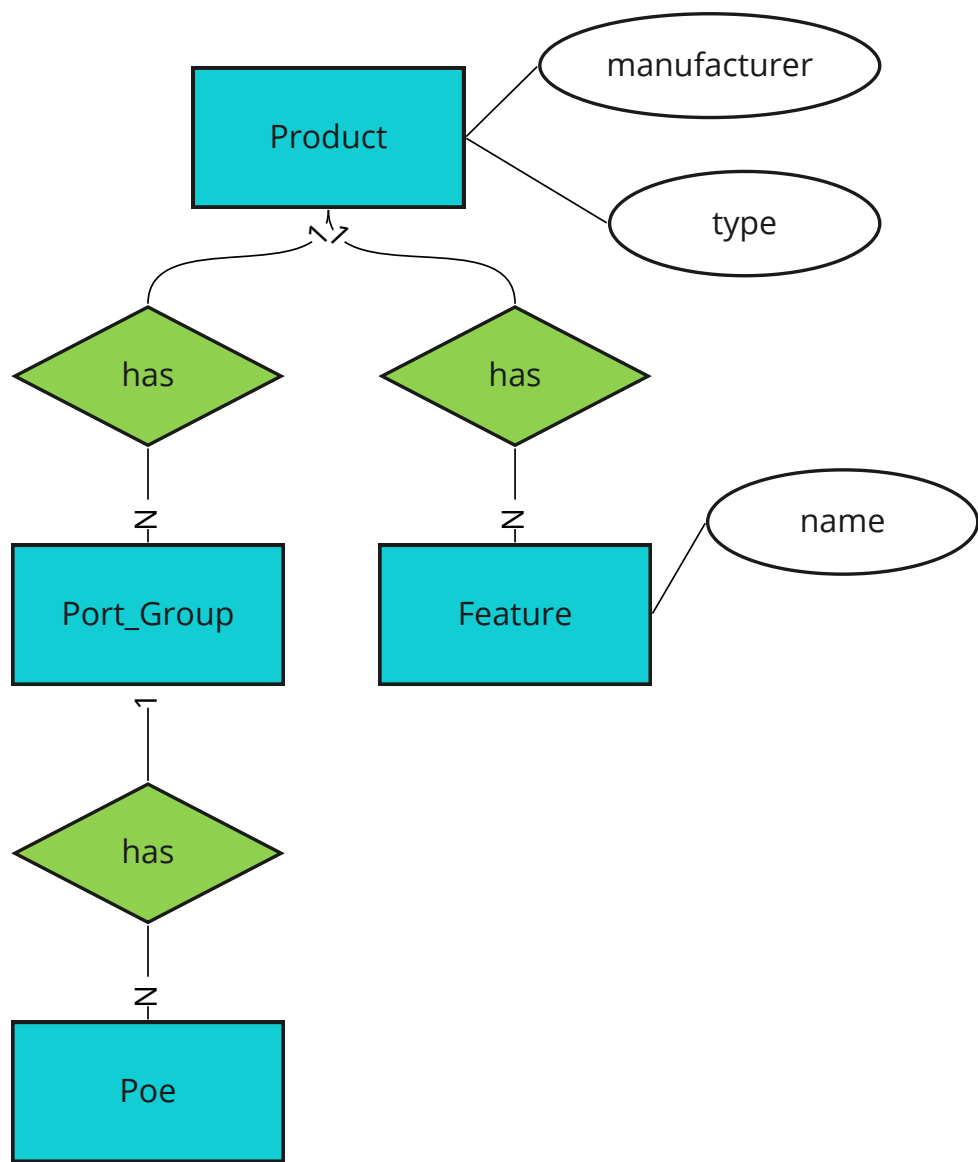
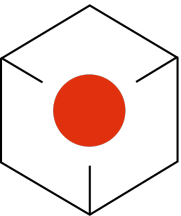
Features können Features enthalten, das ist neu, aber eigentlich auch nur ein weiteres Attribut in Features.

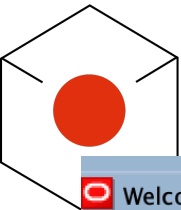
Datenmodell für spezifische Tabellen

Datenmodell für eav Ansatz und weitere Generalisierung indem Port-Group und PoE auch zu Feature wird.

EAV







# Datenstruktur

```
ALTER TABLE feature ADD amount_of_rules NUMBER(10);
ALTER TABLE feature ADD amount_of_connections NUMBER(10);
ALTER TABLE feature ADD amount_of_nat_rules NUMBER(10);
ALTER TABLE feature ADD ssl_inspection VARCHAR2(10);
ALTER TABLE feature ADD amount_of_tunnels NUMBER(10);
ALTER TABLE feature ADD amount_of_users NUMBER(10);
ALTER TABLE feature ADD parent_feat_id NUMBER(10);

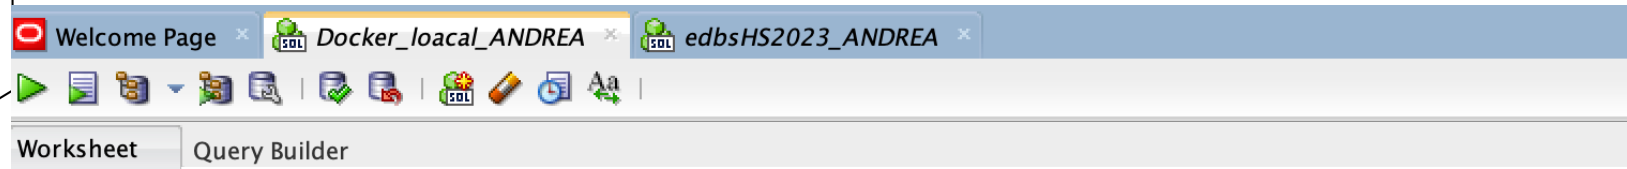
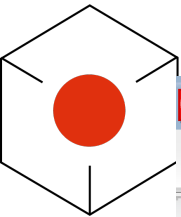
ALTER TABLE feature ADD CONSTRAINT feat_feat_fk
FOREIGN KEY (parent_feat_id) REFERENCES feature (feat_id);
```

```
ALTER TABLE feature_eav ADD parent_feat_id NUMBER(10);
ALTER TABLE feature_eav ADD CONSTRAINT feae_feae_fk
FOREIGN KEY (parent_feat_id) REFERENCES feature_eav (feat_id);
```

EAV

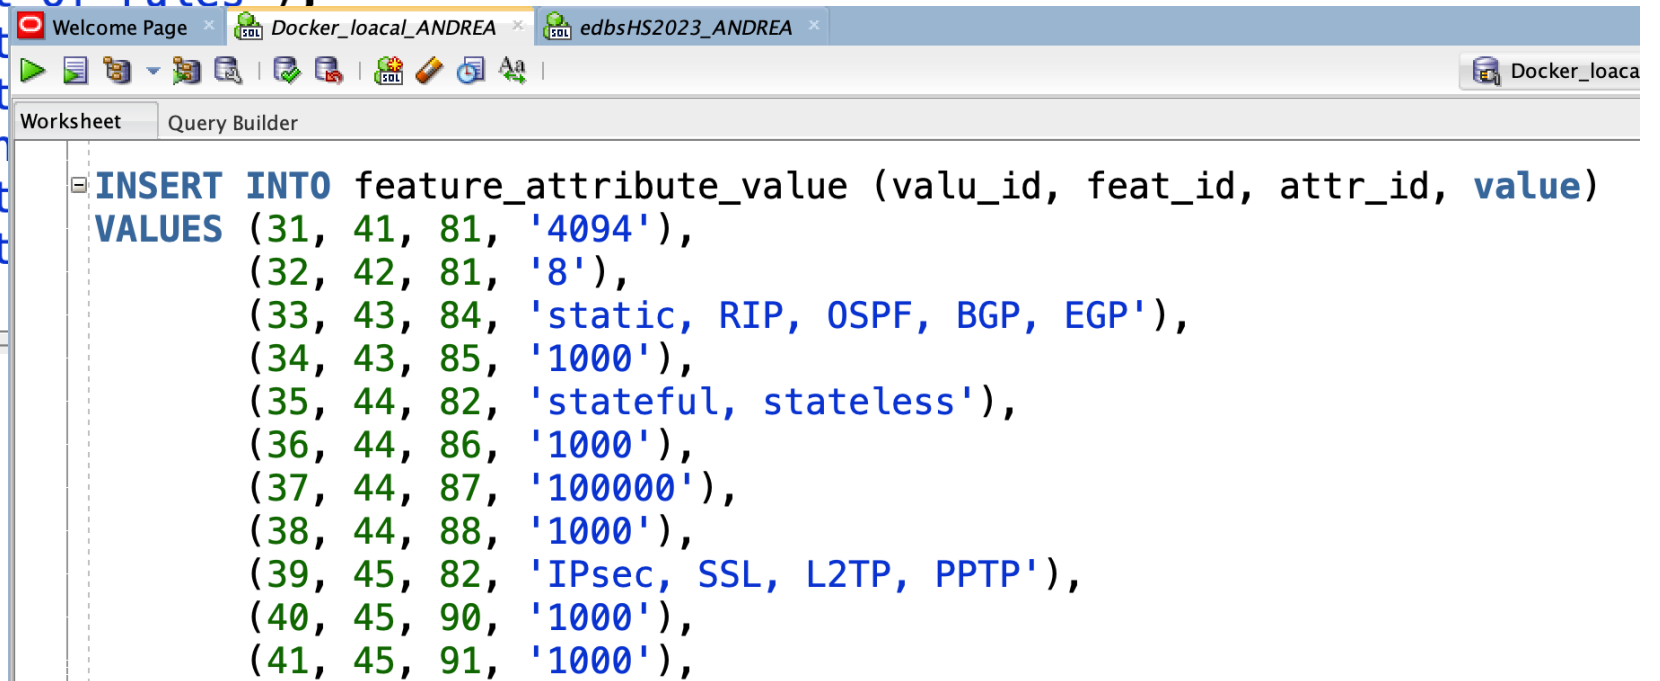




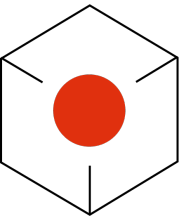


# Neue Daten

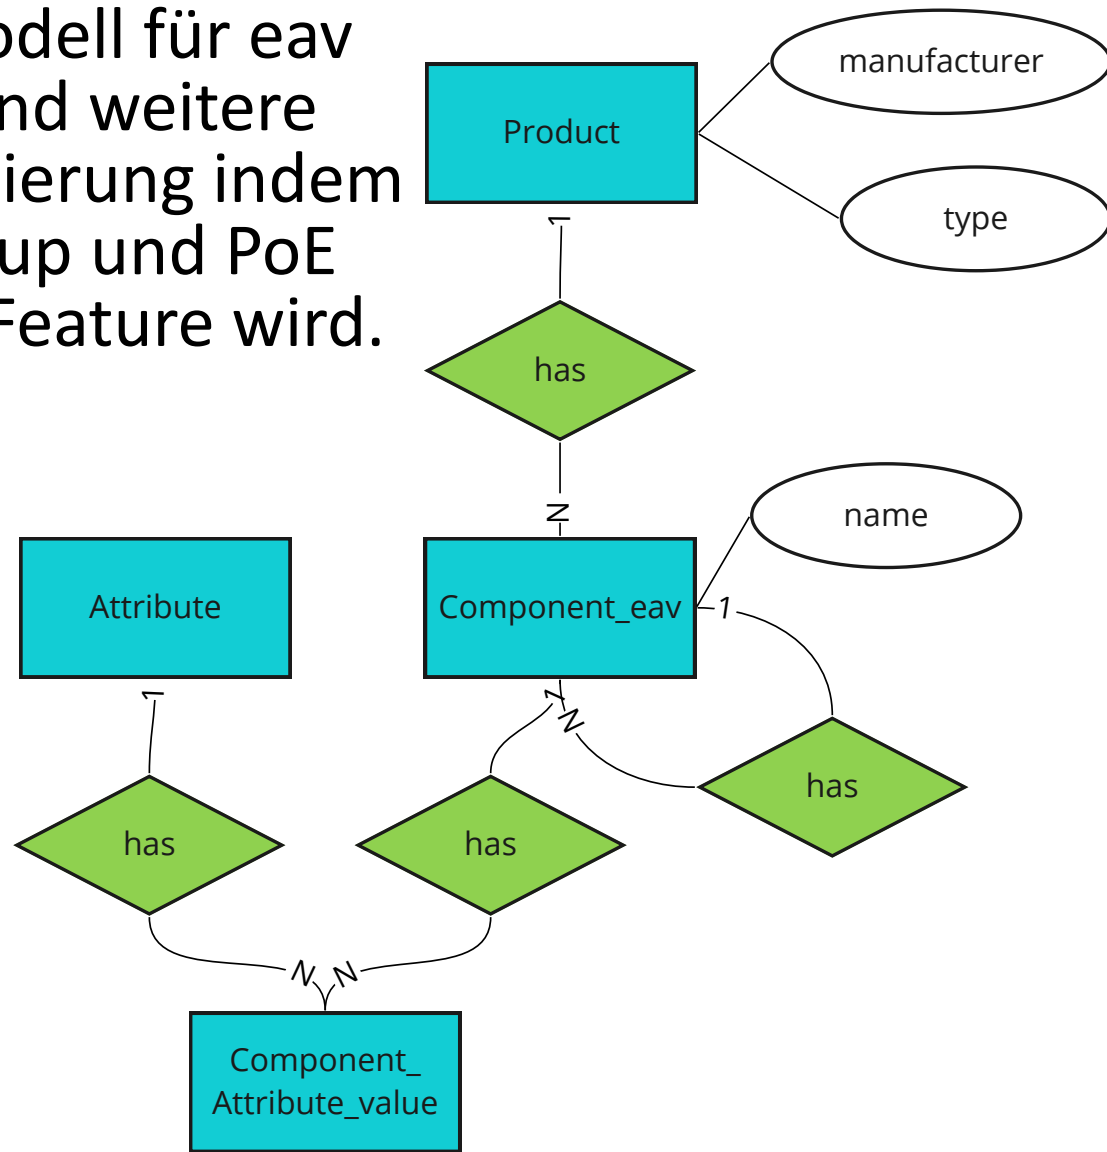
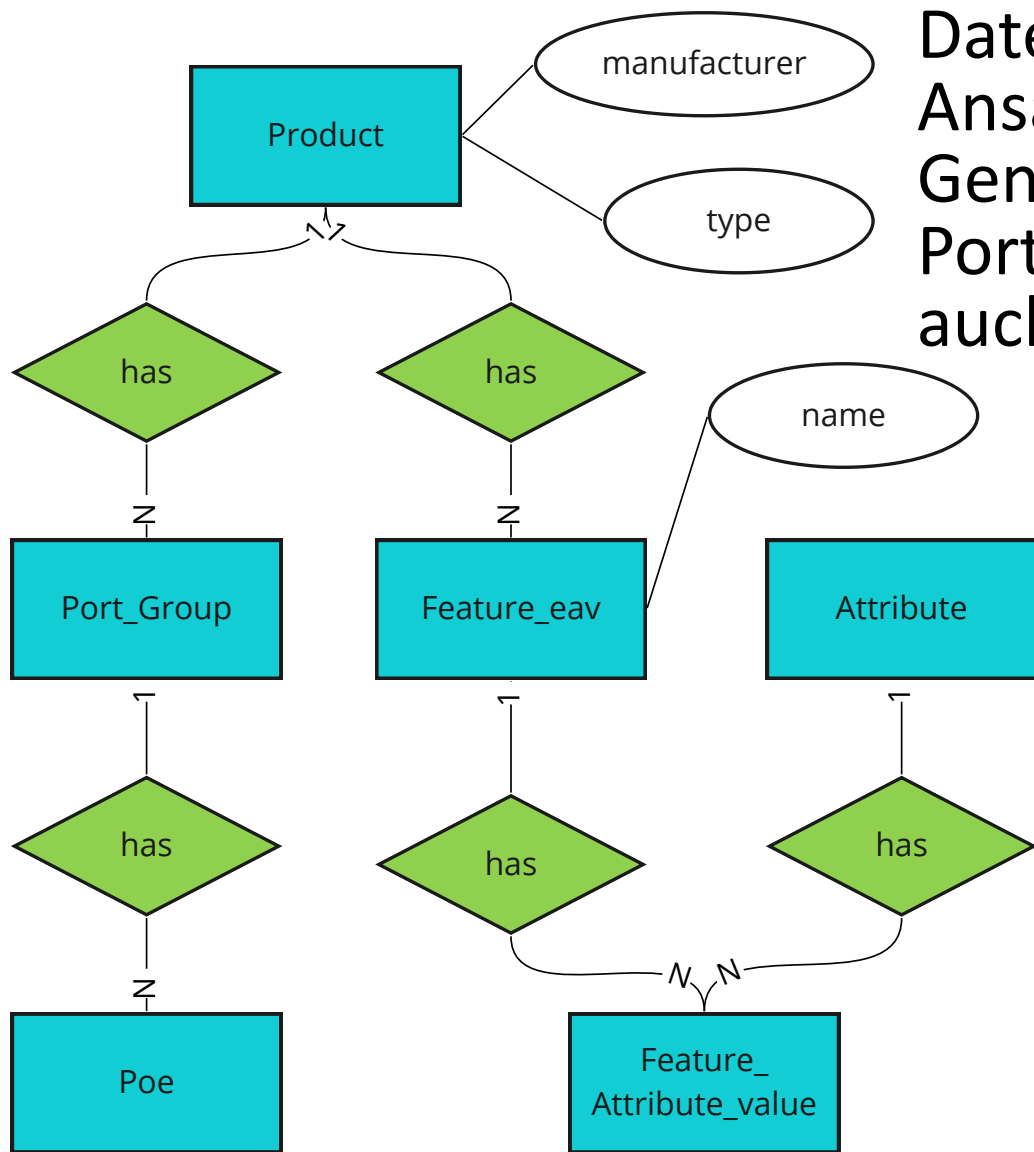
```
+ INSERT INTO feature_eav (feat_id, prod_id, name)...  
+ INSERT INTO feature_eav (feat_id, parent_feat_id, name)...  
- INSERT INTO attribute (attr_id, name)  
VALUES (86, 'amount of rules').  
(87, 'amount...  
(88, 'amount...  
(89, 'ssl_in...  
(90, 'amount...  
(91, 'amount...
```



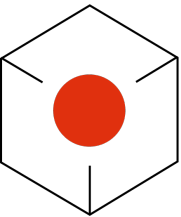
# EAV



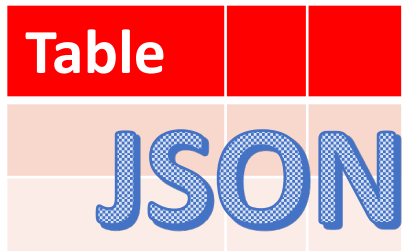
Datenmodell für eav  
Ansatz und weitere  
Generalsierung indem  
Port-Group und PoE  
auch zu Feature wird.



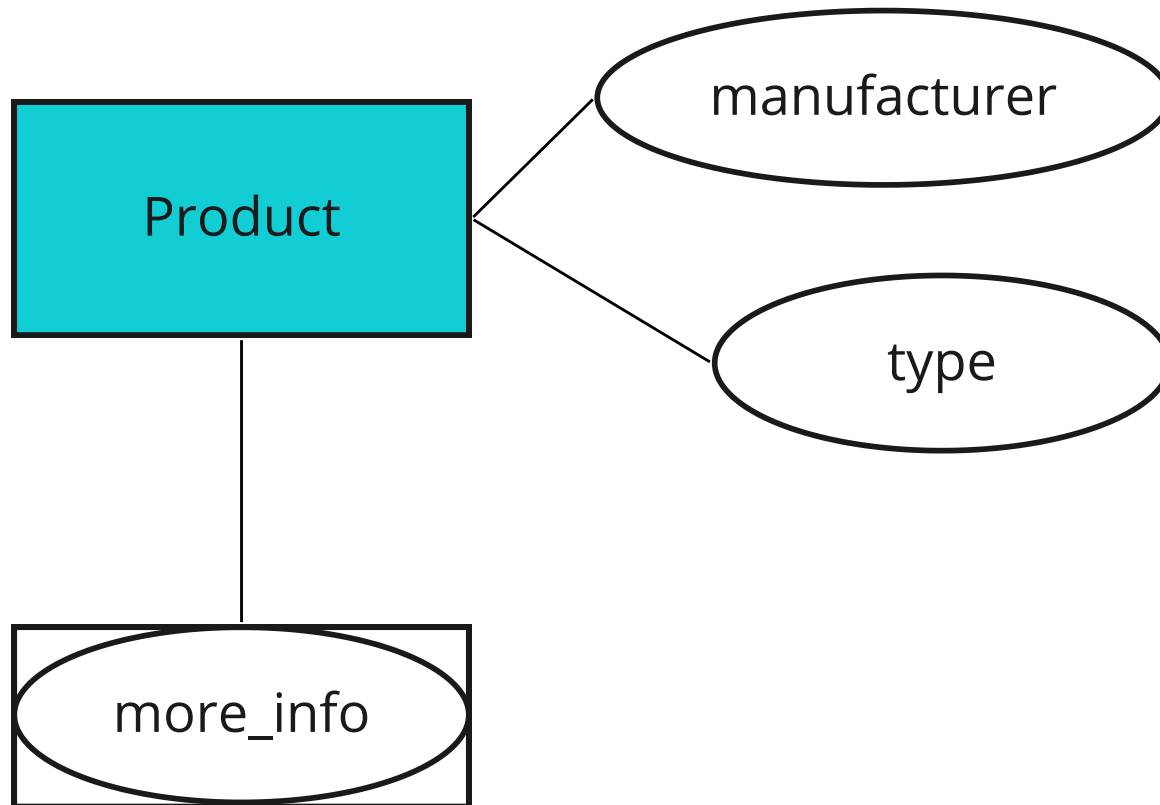
EAV



# Gäbe es da noch andere Lösungen?

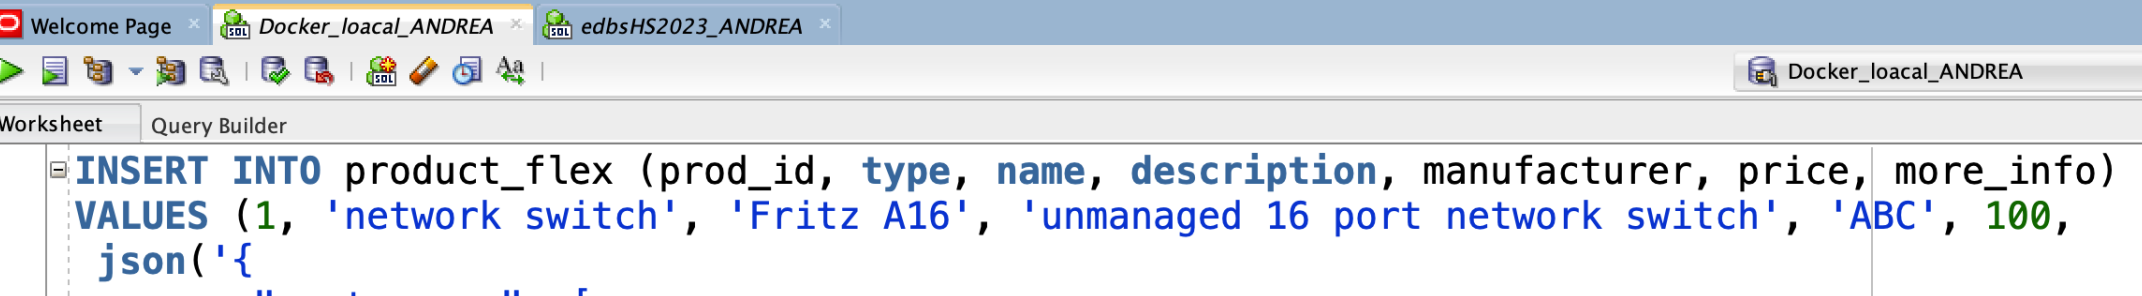


JSON in Attribut  
Prinzip erklären  
Datenmodell  
Demo





# JSON

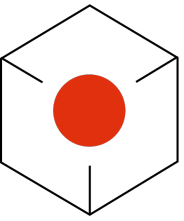


The screenshot shows a SQL IDE with three tabs: 'Welcome Page', 'Docker\_loacal\_ANDREA', and 'edbsHS2023\_ANDREA'. The 'Query Builder' tab is active, displaying a SQL query. The query is an INSERT statement into a table named 'product\_flex'. The columns specified are 'prod\_id', 'type', 'name', 'description', 'manufacturer', 'price', and 'more\_info'. The values provided are (1, 'network switch', 'Fritz A16', 'unmanaged 16 port network switch', 'ABC', 100, and a JSON object). The JSON object represents a port group with an amount of 16, type 'RJ45', and speeds '10/100/1000'.

```

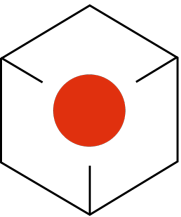
INSERT INTO product_flex (prod_id, type, name, description, manufacturer, price, more_info)
VALUES (1, 'network switch', 'Fritz A16', 'unmanaged 16 port network switch', 'ABC', 100,
json('{
    "port_group": [
        {
            "amount": 16,
            "type": "RJ45",
            "speeds": "10/100/1000"
        }
    ]
}')));

```



# Haben wir Geräte mit einer Port-Group vom Typ SFP?

```
SELECT pf.prod_id, pf.type product_type, pf.name,  
       jt.*  
FROM product_flex pf,  
     JSON_TABLE(  
       pf.more_info  
       COLUMNS (  
         NESTED port_group[*]  
         COLUMNS (  
           amount NUMBER(3) PATH '$.amount',  
           type VARCHAR2(50) PATH '$.type'  
         )  
       )  
     ) jt  
WHERE jt.type = 'SFP';
```



Gib mir alle Angaben zum network switch “Fritz A16”

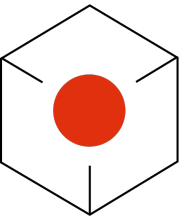
The screenshot shows a SQL query editor with three tabs: 'Welcome Page', 'Docker\_loacal\_ANDREA', and 'edbsHS2023\_ANDREA'. The 'Query Builder' tab is active. It contains three SQL queries:

```
SELECT p.prod_id, p.type product_type, p.name, p.price
FROM product p INNER JOIN ...

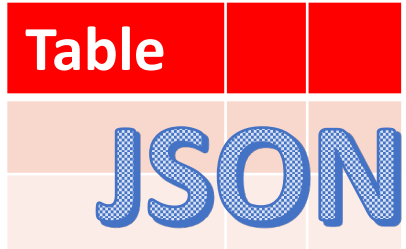
SELECT pc.*
FROM product_collection pc
WHERE pc.json_data.type = 'network switch'
      AND pc.json_data.name = 'Fritz A16';

SELECT pf.*
FROM product_flex pf
WHERE pf.type = 'network switch'
      AND pf.name = 'Fritz A16';
```

Problem,  
wenn JSON ein  
typo hat



# Zusammenfassung und Abstimmung



- Tabelle mit JSON

- EAV

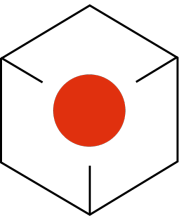
EAV

- Dokument Datenbank



- Relationale Datenbank





# Dr. Andrea Kennel

**SYM<sup>L2</sup>**



Consultant

Dozentin für Datenbanken

Coach für Project Management

Fachhochschule Nordwestschweiz

Brugg/Windisch, Schweiz



[andrea.kennel@fhnw.ch](mailto:andrea.kennel@fhnw.ch)

[andrea@infokennel.ch](mailto:andrea@infokennel.ch)

[www.infokennel.ch](http://www.infokennel.ch)