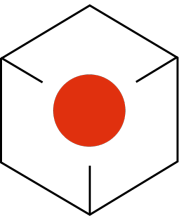




# Dynamic attributes to store flexible data

Dr. Andrea Kennel InfoPunkt Kennel GmbH Switzerland



# Dr. Andrea Kennel



Consultant

Lecturer for Databases

Coach for Project Management

University of Applied Sciences  
Northwestern Switzerland

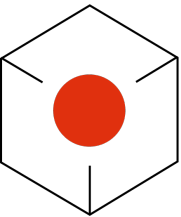
Brugg/Windisch, Switzerland



[andrea.kennel@fhnw.ch](mailto:andrea.kennel@fhnw.ch)

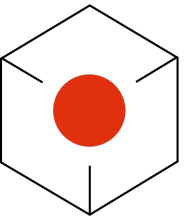
[andrea@infokennel.ch](mailto:andrea@infokennel.ch)

[www.infokennel.ch](http://www.infokennel.ch)



```
{  "type": "router",
  "name": "Hans B4",
  "description": "4 port router",
  "manufacturer": "XYZ",
  "price": 800,
  "port_group": [
    { "amount": 4,
      "type": "RJ45",
      "speeds": "100/1000"
    },
    { "amount": 2,
      "type": "SFP",
      "speeds": "1000/10000"
    }
  ],
  "routing": {
    "protocols": "static OSPF",
    "table_size": 5
  }
}
```

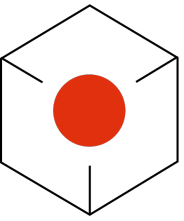




# How do the devices look?

```
{  "type": "router",
    "name": "Hans B4",
    "description": "4 port router",
    "manufacturer": "XYZ",
    "price": 800,
    "port_group": [
      { "amount": 4,
        "type": "RJ45",
        "speeds": "100/1000"
      },
```

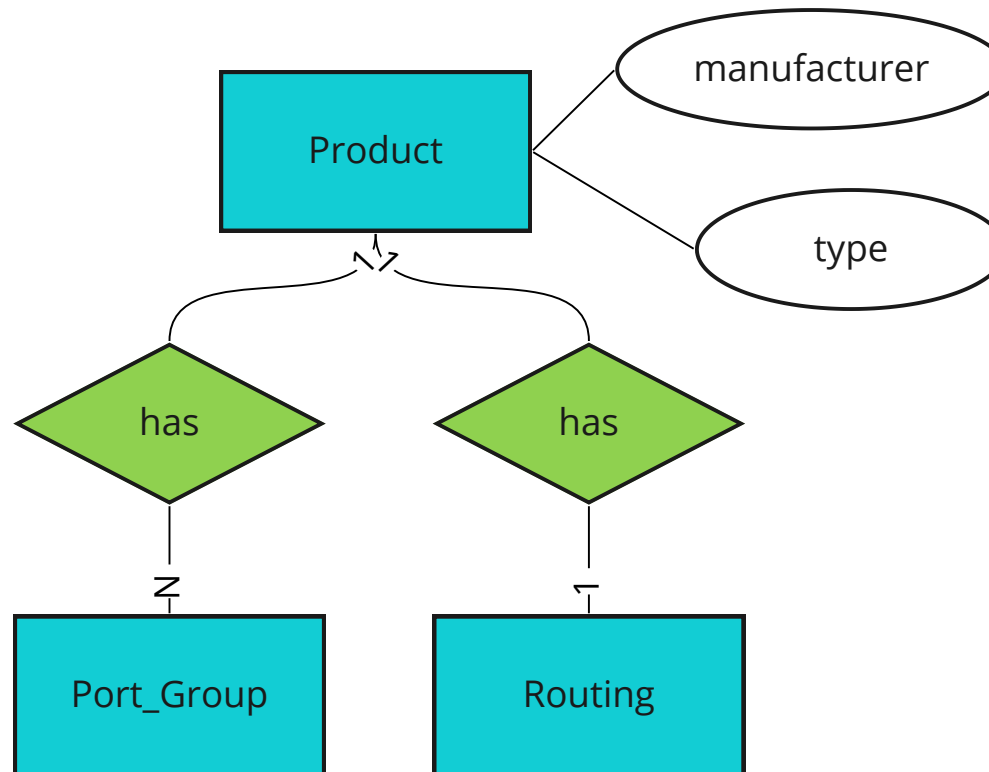
```
    { "amount": 2,
      "type": "SFP",
      "speeds": "1000/10000"
    },
    "routing": {
      "protocols": "static OSPF",
      "table_size": 5
    }
}
```




# Could you quickly ...

*For a web shop that sells network devices, the data on the network devices must be stored in a database.*

*The data is  
as JSON,  
the model is  
relatively simple:*





# Implement data model and enter data

```
CREATE TABLE product (  
  prod_id      NUMBER(10),  
  type         VARCHAR2(50),  
  name         VARCHAR2(50),  
  description   VARCHAR2(200),  
  manufacturer  VARCHAR2(50),  
  price        NUMBER (8,2)  
);  
  
CREATE TABLE port_group (  

```

Query Result x Script Output x  
Task completed in 1.972 seconds

Table PRODUCT created.

Table PORT\_GROUP created.

Table ROUTING created.

```
INSERT INTO product (prod_id, type, name, description, manufacturer  
VALUES (1, 'network switch', 'Fritz A16', 'unmanaged 16 port network  
(2, 'network switch', 'Fritz B24', 'unmanaged 24 port 10gbit  
(3, 'network switch', 'Fritz C12', 'managed 12 port network  
(4, 'router', 'Hans A8', '8 port router', 'ABC', 1000),  
(5, 'router', 'Hans B4', '4 port router', 'XYZ', 800);  
  
INSERT INTO port_group (pogr_id, prod_id, amount, type, speeds)...  
  
INSERT INTO routing (rout_id, prod_id, protocols, table_size)
```

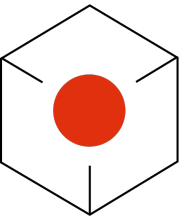
Query Result x Script Output x  
Task completed in 0.597 seconds

5 rows inserted.

8 rows inserted.

2 rows inserted.

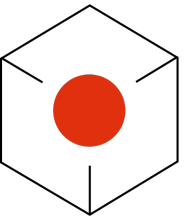




# We have a new device

```
{ "type": "network switch/layer3 switch",  
  "name": "Fritz C16",  
  "description": "16 port PoE layer 3 network switch",  
  "manufacturer": "ABC",  
  "price": 500,  
  "port_group": [  
    { "amount": 16,  
      "type": "RJ45",  
      "speeds": "10/100/1000",  
      "poe": {"modes": ["active", "passive"],  
              "volt": [24, 48]}  
    },  
  ],  
  ...
```

```
    "feature": [  
      { "name": "VLAN",  
        "amount": 4094},  
      { "name": "QoS",  
        "amount": 8},  
      { "name": "network access control",  
        "type": "MAC based authentication",  
        "vlan_support": true},  
      { "name": "routing",  
        "protocols": "static, RIP, OSPF, BGP",  
        "table_size": 10}  
    ],  
  }
```



# Then we have to adapt the data structure

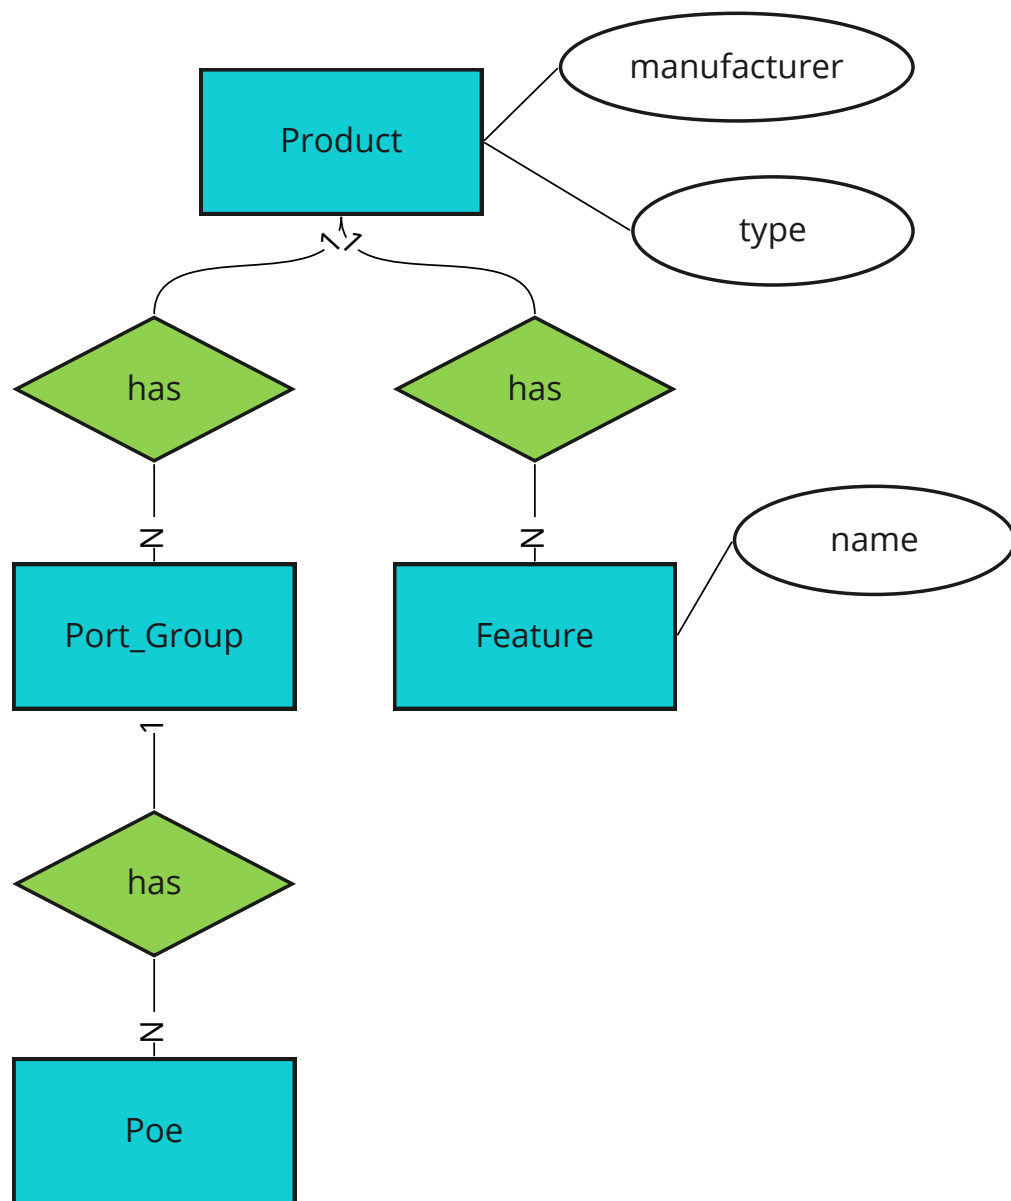
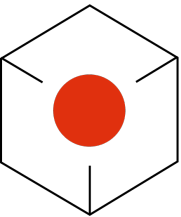
Detail table for port group and

Feature as generalisation with several specialisations

We don't know what else is coming and combine all attributes in the generalisation.

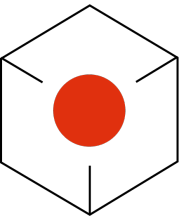






# New data structure





# New data structure

Worksheet | Query Builder

```
CREATE TABLE poe (  
  poe_id      NUMBER(10),  
  pogr_id     NUMBER(10),  
  modes       VARCHAR2(20),  
  volt        NUMBER (6,2)  
);  
  
ALTER TABLE poe ADD CONSTRAINT poe_pk PRIMARY KEY (poe_id);  
ALTER TABLE poe ADD CONSTRAINT poe_pogr_fk FOREIGN KEY (pogr_id)  
REFERENCES port_group (pogr_id);
```

Query Result | Script Output


Task completed in 0.555 seconds

Table P0E created.

Table P0E altered.

Table P0E altered.





# data structure

Worksheet | Query Builder

```
RENAME routing TO feature;  
ALTER TABLE feature RENAME COLUMN rout_id TO feat_id;  
ALTER TABLE feature ADD name VARCHAR2(200);  
ALTER TABLE feature ADD amount NUMBER(10);  
ALTER TABLE feature ADD type VARCHAR2(200);  
ALTER TABLE feature ADD vlan_support VARCHAR2(200);  
-- all existing get name 'routing'  
UPDATE feature  
SET name = 'routing';
```

Query Result | Script Output

Task completed in 1.832 seconds

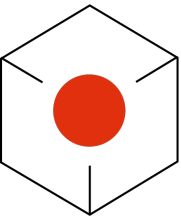
Table FEATURE altered.

Table FEATURE altered.

2 rows updated.







# What if there are more attributes?

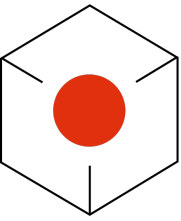
We had to add more attributes to FEATURES.

What happens if there are more attributes?

What if many features only have a few attributes filled?

Is there a more flexible solution?

```
"feature": [  
  { "name": "VLAN",  
    "amount": 4094},  
  { "name": "QoS",  
    "amount": 8},  
  { "name": "network access control",  
    "type": "MAC based authentication",  
    "vlan_support": true},  
  { "name": "routing",  
    "protocols": "static, RIP, OSPF, BGP",  
    "table_size": 10}  
]  
}
```



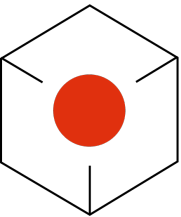
# What if there are more attributes?

The answer is EAV

The Entity Attribute Value (EAV) model is a data modelling technique used in databases to store and retrieve data in a flexible and scalable way.

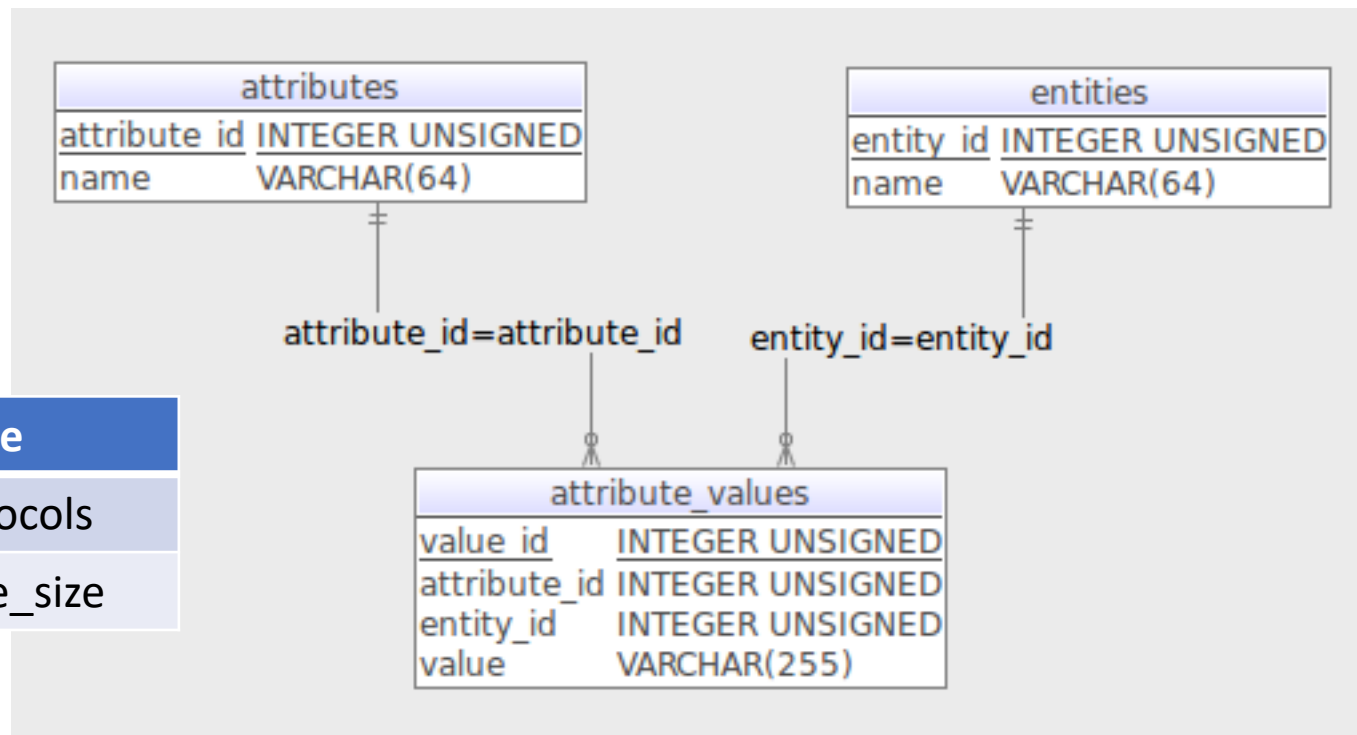
Details about the model: <https://inviqa.com/blog/understanding-eav-data-model-and-when-use-it>

**EAV**



# EAV

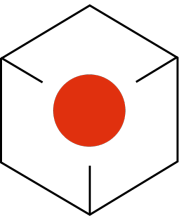
Attribute_id	name
84	protocols
85	table_size



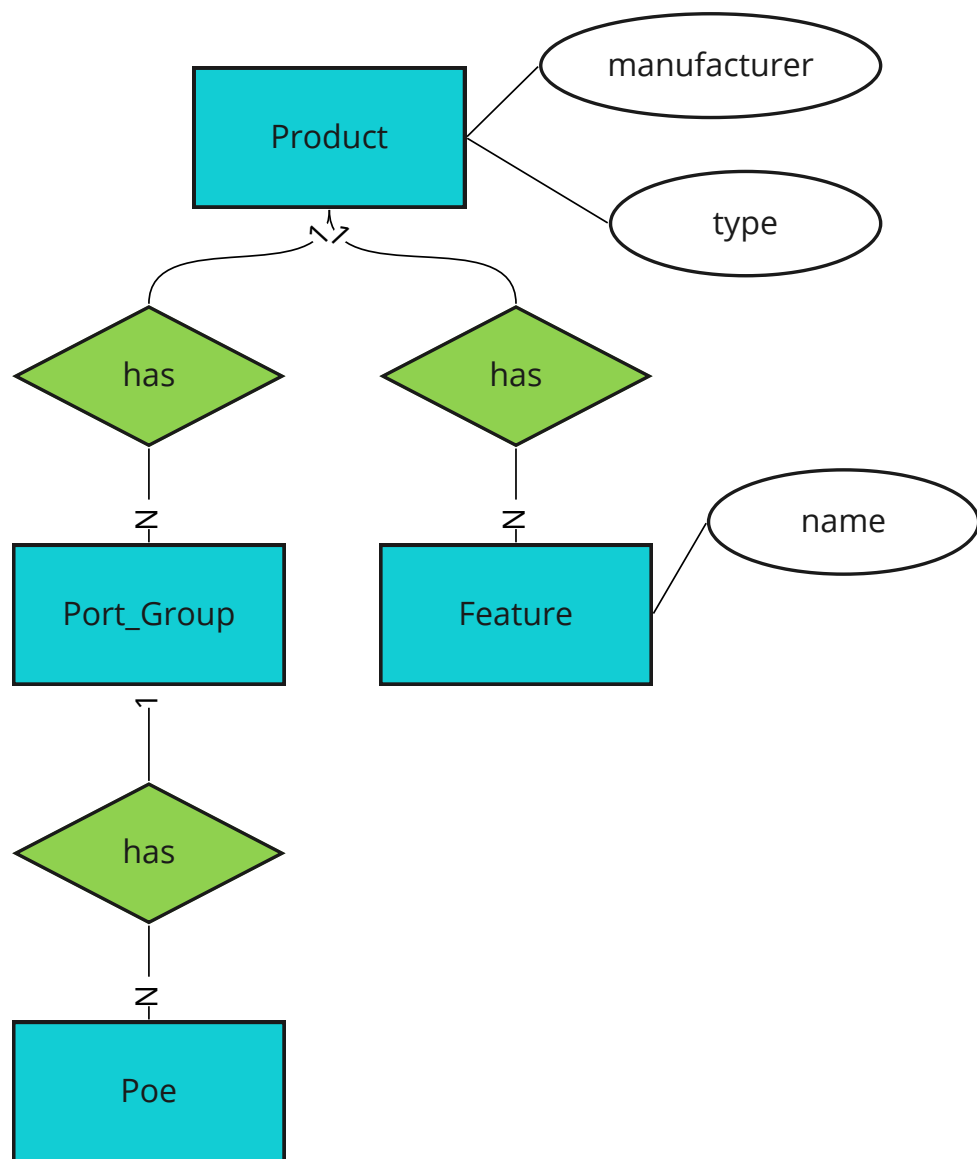
```
INSERT INTO feature (feat_id, prod_id, name, protocols, table_size)
VALUES (24, 21, 'routing', 'static, RIP, OSPF, BGP', 10);
```

Value_id	Attribute_id	Entity_id	values
1	84	24	static, RIP, OSPF, BGP
2	85	24	10

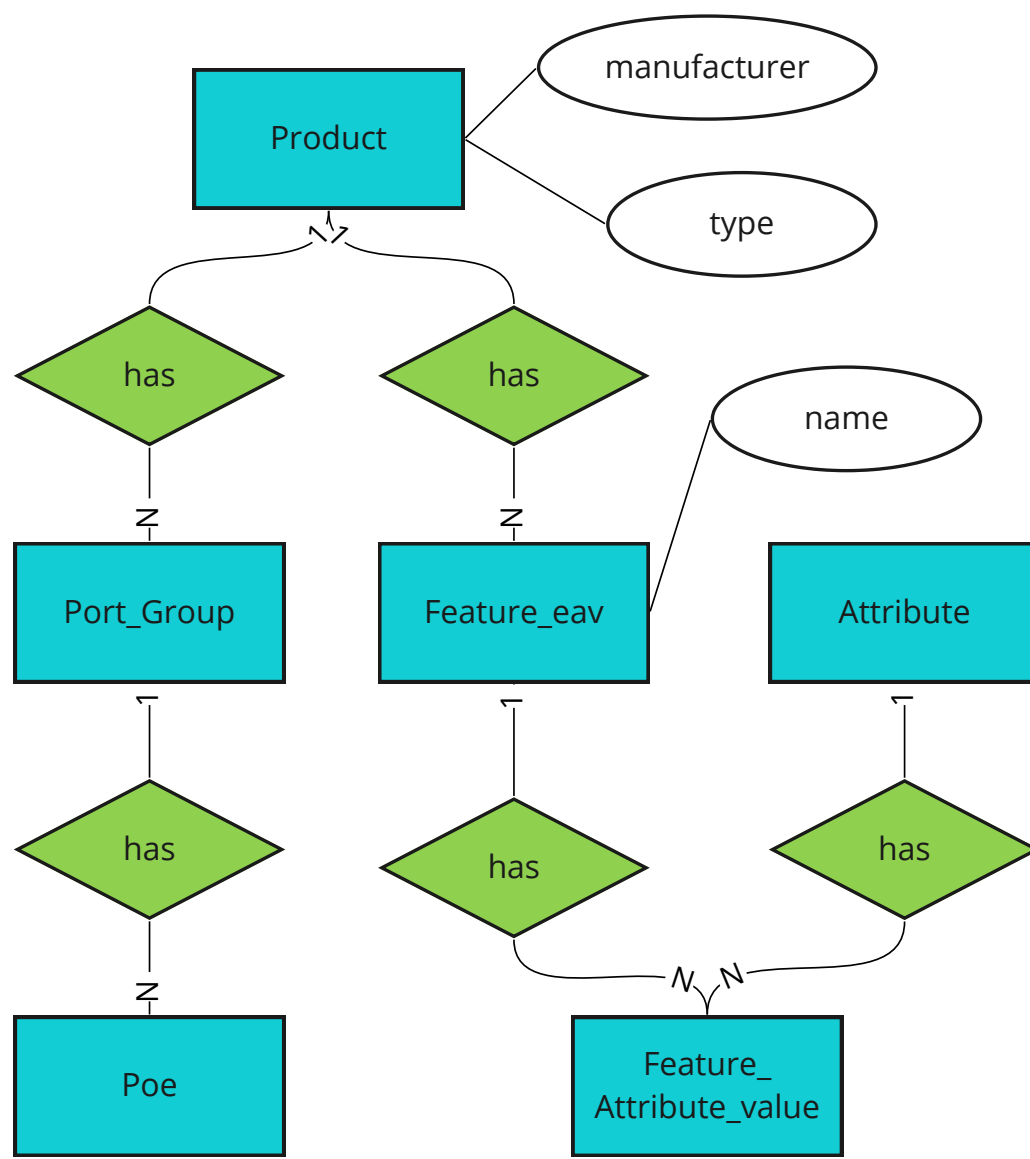
# EAV



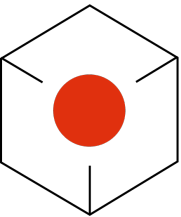
EAV



EAV







# EAV

```
Welcome Page | Docker_loacal_ANDREA | edbsHS2023_AN

Worksheet | Query Builder

CREATE TABLE feature_eav (
  feat_id      NUMBER(10),
  prod_id      NUMBER(10),
  name         VARCHAR2(200)
);

CREATE TABLE attribute (
  attr_id      NUMBER(10),
  name         VARCHAR2(30)
);

CREATE TABLE feature_attribute_value (
  valu_id      NUMBER(10),
  feat_id      NUMBER(10),
  attr_id      NUMBER(10),
  value        VARCHAR2(2000)
);
```

```
Welcome Page | Docker_loacal_ANDREA | edbsHS2023_ANDREA

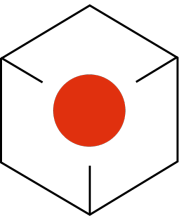
Worksheet | Query Builder

INSERT INTO feature_eav (feat_id, prod_id, name)
VALUES (24, 21, 'routing');

INSERT INTO attribute (attr_id, name)
VALUES (84, 'protocols'),
       (85, 'table_size');

INSERT INTO feature_attribute_value (valu_id, feat_id, attr_id, value)
VALUES (25, 24, 84, 'static, RIP, OSPF, BGP'),
       (26, 24, 85, '10');
```

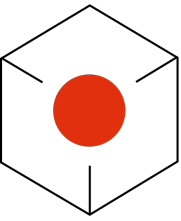
# EAV



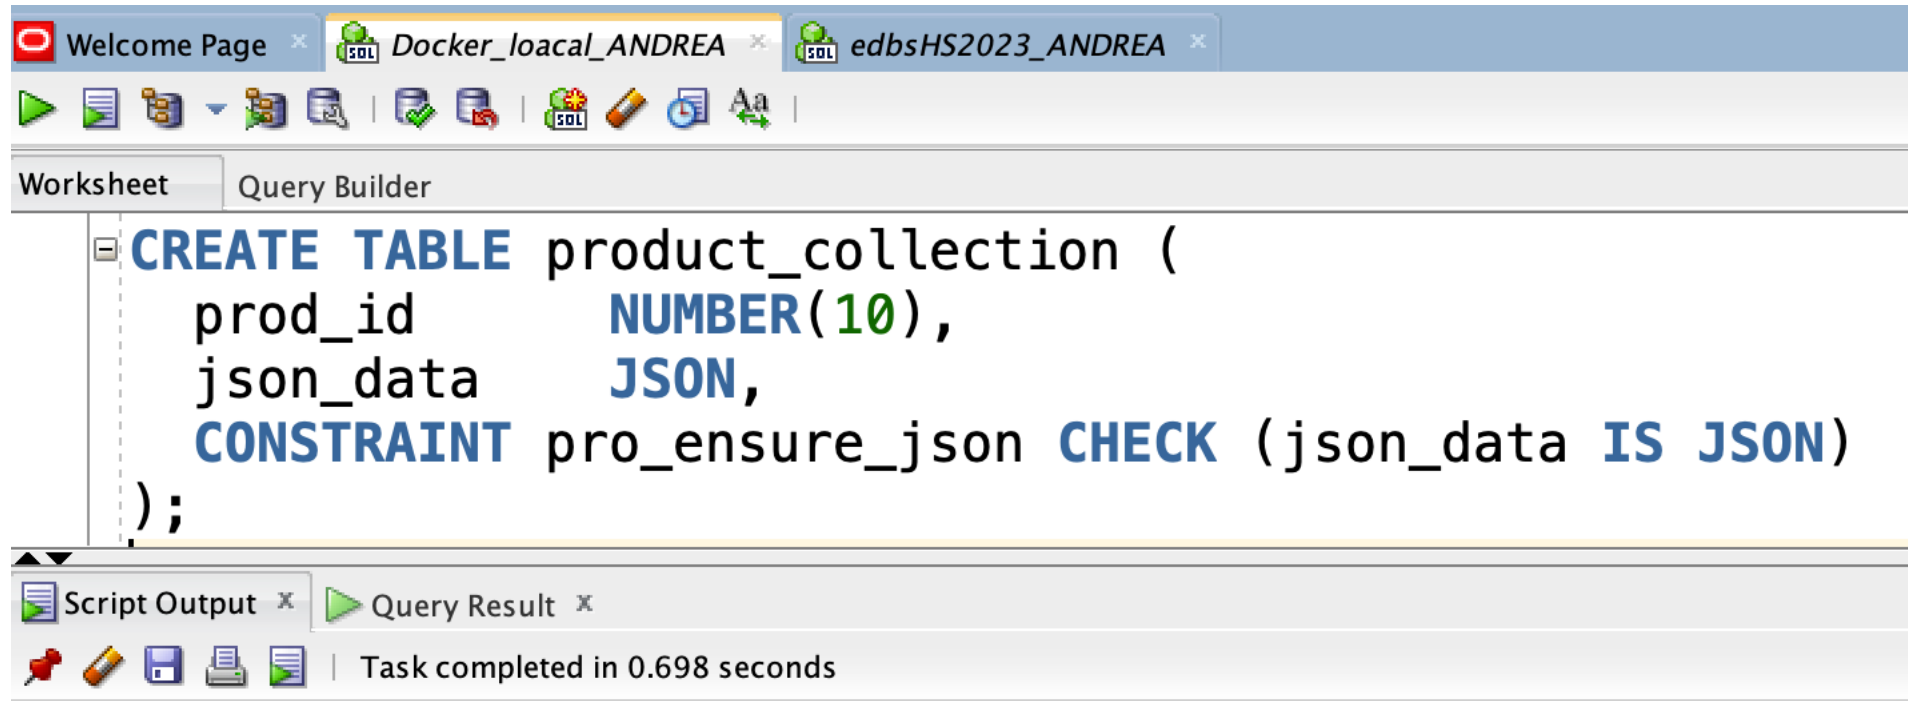
# Wouldn't a document DB be better?

- I am getting confused
- Why not just store all the data as JSON?





# Wouldn't a document DB be better?



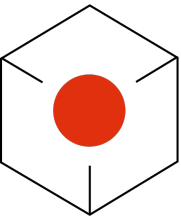
The screenshot shows a SQL IDE window with three tabs: 'Welcome Page', 'Docker\_loacal\_ANDREA', and 'edbsHS2023\_ANDREA'. The 'Query Builder' tab is active, displaying the following SQL code:

```
CREATE TABLE product_collection (  
  prod_id      NUMBER(10),  
  json_data    JSON,  
  CONSTRAINT pro_ensure_json CHECK (json_data IS JSON)  
);
```

Below the code editor, the 'Script Output' and 'Query Result' tabs are visible. The status bar at the bottom indicates 'Task completed in 0.698 seconds'.



Table PRODUCT\_COLLECTION created.



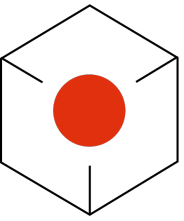
# Wouldn't a document DB be better?

The screenshot displays two side-by-side windows of a SQL IDE. Both windows have tabs for 'Welcome Page', 'Docker\_loacal\_ANDREA', and 'edbsHS2023\_ANDREA'. The left window is in 'Worksheet' mode and shows a SQL query:

```
INSERT INTO product_collection (prod_id, json_data)
VALUES (1, json('{
  "type": "network switch",
  "name": "Fritz A16",
  "description": "unmanaged 16 port network switch",
  "manufacturer": "ABC",
  "price": 100,
  "port_group": [
    {
      "amount": 16,
      "type": "RJ45",
      "speeds": "10/100/1000"
    }
  ]
}'));
```

The right window is in 'Query Builder' mode and shows a more detailed JSON document for a Fritz C16 switch:

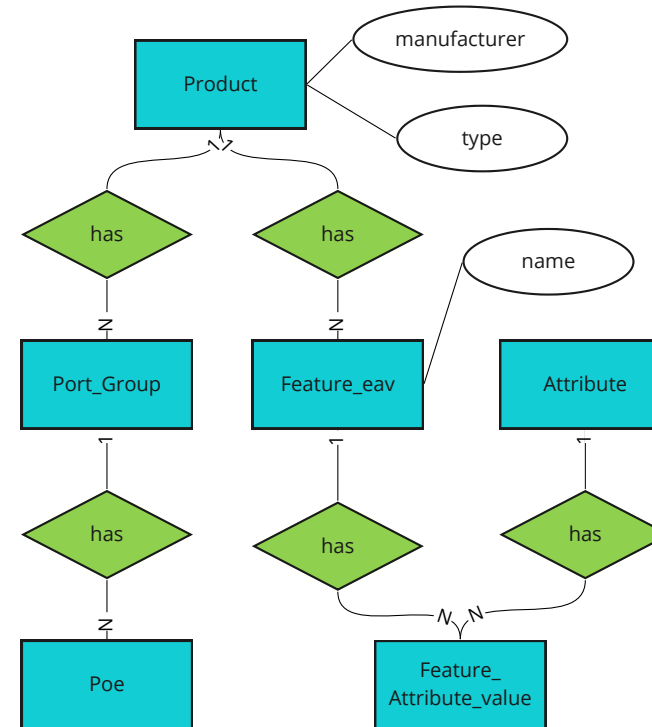
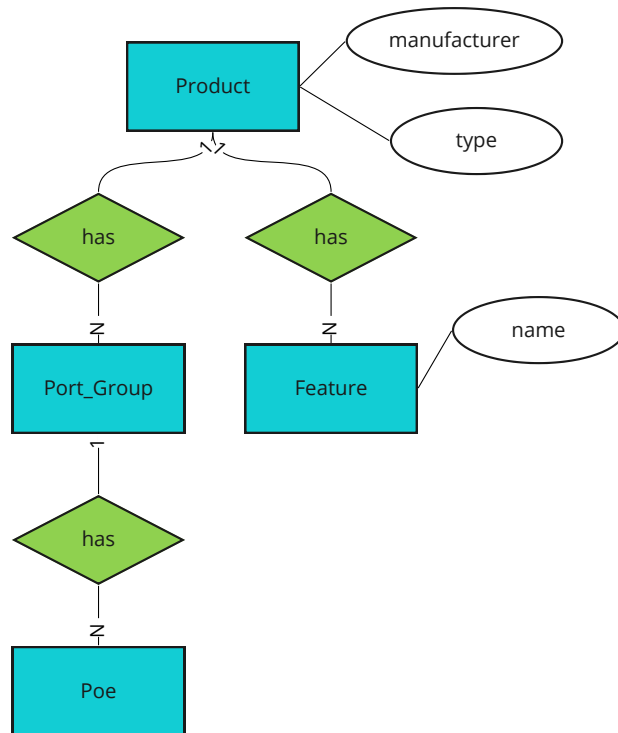
```
INSERT INTO product_collection (prod_id, json_data)
VALUES (21, json('{
  "type": "network switch/layer3 switch",
  "name": "Fritz C16",
  "description": "16 port PoE layer 3 network switch",
  "manufacturer": "ABC",
  "price": 500,
  "port_group": [
    {
      "amount": 16,
      "type": "RJ45",
      "speeds": "10/100/1000",
      "poe": {
        "modes": [
          "active",
          "passive"
        ],
        "volt": [
          24,
          48
        ]
      }
    }
  ],
  "feature": [
    {
      "name": "VLAN",
      "amount": 4094
    },
    {
      "name": "QoS",
      "amount": 8
    },
    {
      "name": "network access control",
      "type": "MAC based authentication",
      "vlan_support": true
    },
    {
      "name": "routing",
      "protocols": "static, RIP, OSPF, BGP",
      "table_size": 10
    }
  ]
}'));
```

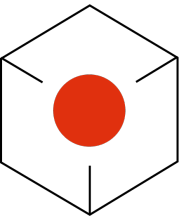


# Wouldn't a document DB be better?

What does it look like with queries?

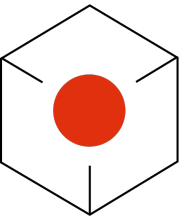
Example: Do we have devices with a port group of type SFP?





Do we have devices with a port group of type SFP?  
relational

```
SELECT pro.prod_id, pro.type product_type, pro.name, pog.amount, pog.type
FROM product pro INNER JOIN
    port_group pog ON (pro.prod_id = pog.prod_id)
WHERE pog.type = 'SFP';
```

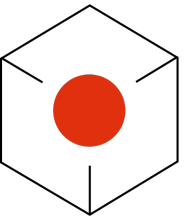


# Do we have devices with a port group of type SFP?

## JSON

```
SELECT pc.prod_id, pc.json_data.type product_type,  
       pc.json_data.name,  
       pc.json_data.port_group[*].amount,  
       pc.json_data.port_group[*].type  
FROM product_collection pc  
WHERE json_exists(pc.json_data,  
                  '$.port_group?(@.type == $v1)'  
                  PASSING 'SFP' AS "v1");
```

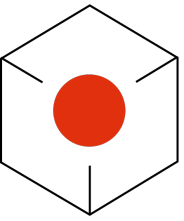
```
SELECT pc.prod_id, pc.json_data.type product_type,  
       pc.json_data.name,  
       jt.*  
FROM product_collection pc,  
     JSON_TABLE(  
       pc.json_data  
       COLUMNS (  
         NESTED port_group[*]  
         COLUMNS (  
           amount NUMBER(3) PATH '$.amount',  
           type VARCHAR2(50) PATH '$.type'  
         )  
       )  
     ) jt  
WHERE jt.type = 'SFP';
```



Do we have feature devices that have amount = 8?  
relational

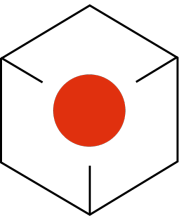
```
SELECT pro.prod_id, pro.type product_type, pro.name,  
       fe.name feature_name, fe.amount  
FROM product pro INNER JOIN  
       feature fe ON (pro.prod_id = fe.prod_id)  
WHERE fe.amount = 8;
```





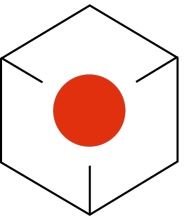
# Do we have feature devices that have amount = 8? JSON

```
SELECT pc.prod_id, pc.json_data.type product_type, pc.json_data.name,  
       jt.*  
FROM product_collection pc,  
     JSON_TABLE(  
       pc.json_data  
       COLUMNS (  
         NESTED feature[*]  
         COLUMNS (  
           feature_name VARCHAR2(50) PATH '$.name',  
           amount NUMBER(10) PATH '$.amount'  
         )  
       )  
     ) jt  
WHERE jt.amount = 8;
```




Do we have feature devices that have amount = 8?  
EAV


```
SELECT pro.prod_id, pro.type product_type, pro.name,  
       fe.name feature_name, feat.value amount  
FROM product pro INNER JOIN  
     feature_eav fe ON (pro.prod_id = fe.prod_id) INNER JOIN  
     feature_attribute_value feat ON (fe.feat_id = feat.feat_id) INNER JOIN  
     attribute at ON (feat.attr_id = at.attr_id)  
WHERE at.name = 'amount'  
      AND feat.value = '8';
```



We have a new



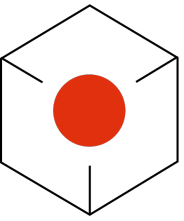
<https://makeameme.org/meme/oh-no-not-6bc4c7>



# We have a new device again

```
{ "type": "router/firewall",  
  "name": "Hans A48",  
  ....  
  "feature": [  
    { ... },  
    { ... },  
    { ... },  
    { "name": "firewall",  
      "type": "stateful, stateless",  
      "amount_of_rules": 1000,  
      "amount_of_connections": 100000,  
      "amount_of_nat_rules": 1000,  
      "feature": [  
        { "name": "URL filtering",  
          "types": "whitelist, blacklist",  
          "amount_of_rules": 1000 },  
        ....  
      ]  
    }  
  ]  
}
```

```
{ "name": "application filtering",  
  "types": "whitelist",  
  "amount_of_rules": 1000 },  
{ "name": "content filtering",  
  "types": "blacklist",  
  "amount_of_rules": 1000 },  
{ "name": "anti-virus",  
  "ssl_inspection": true }  
]  
{ "name": "VPN",  
  "type": "IPsec, SSL, L2TP, PPTP",  
  "amount_of_tunnels": 1000,  
  "amount_of_users": 1000,  
  "feature": [  
    { "name": "IPsec",  
      "type": "IKEv1, IKEv2",  
      "amount_of_tunnels": 1000,  
      "amount_of_users": 1000 },  
    { "name": "SSL",  
      ....  
    }  
  ]  
}
```



# Then we have to adapt the data structure

Further types of features with new attributes, we have already expected that.

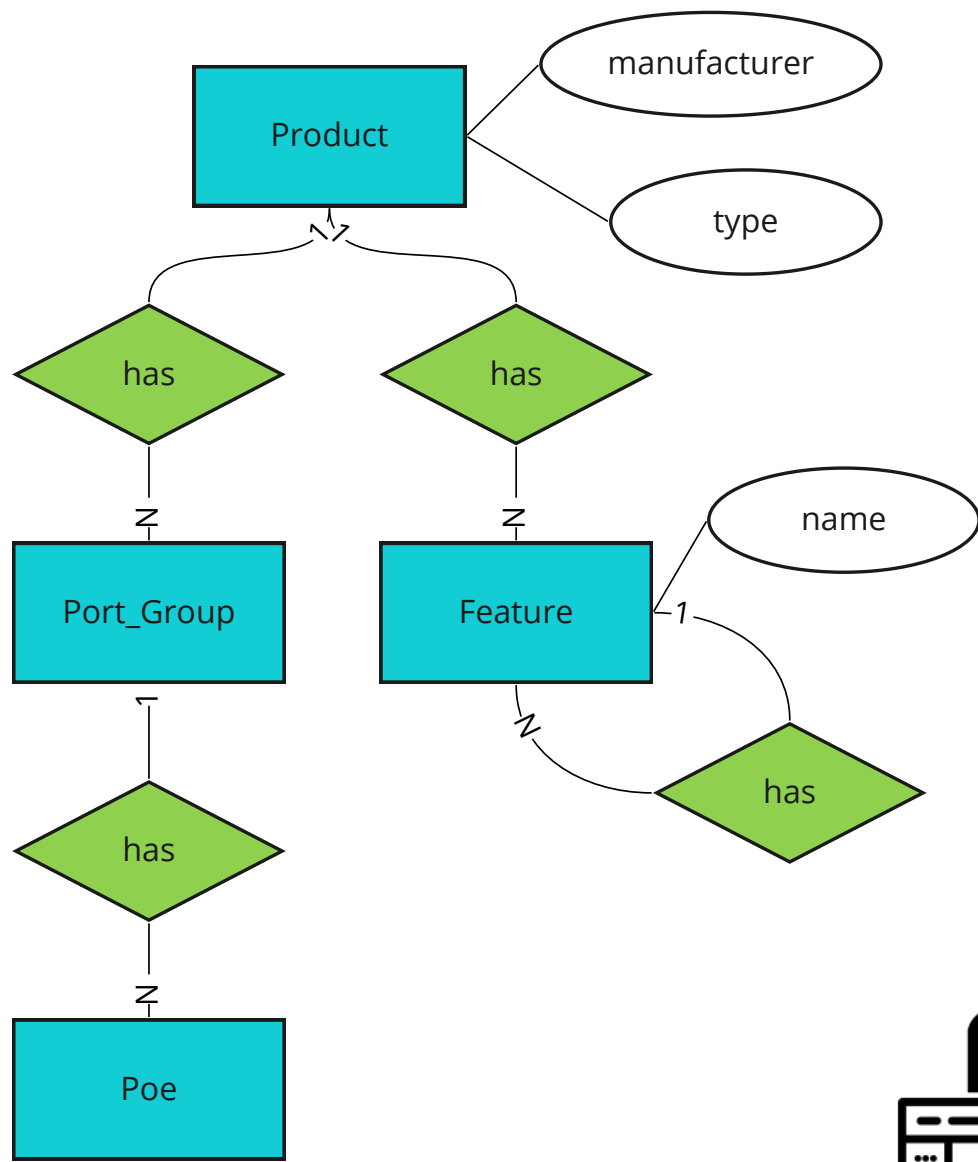
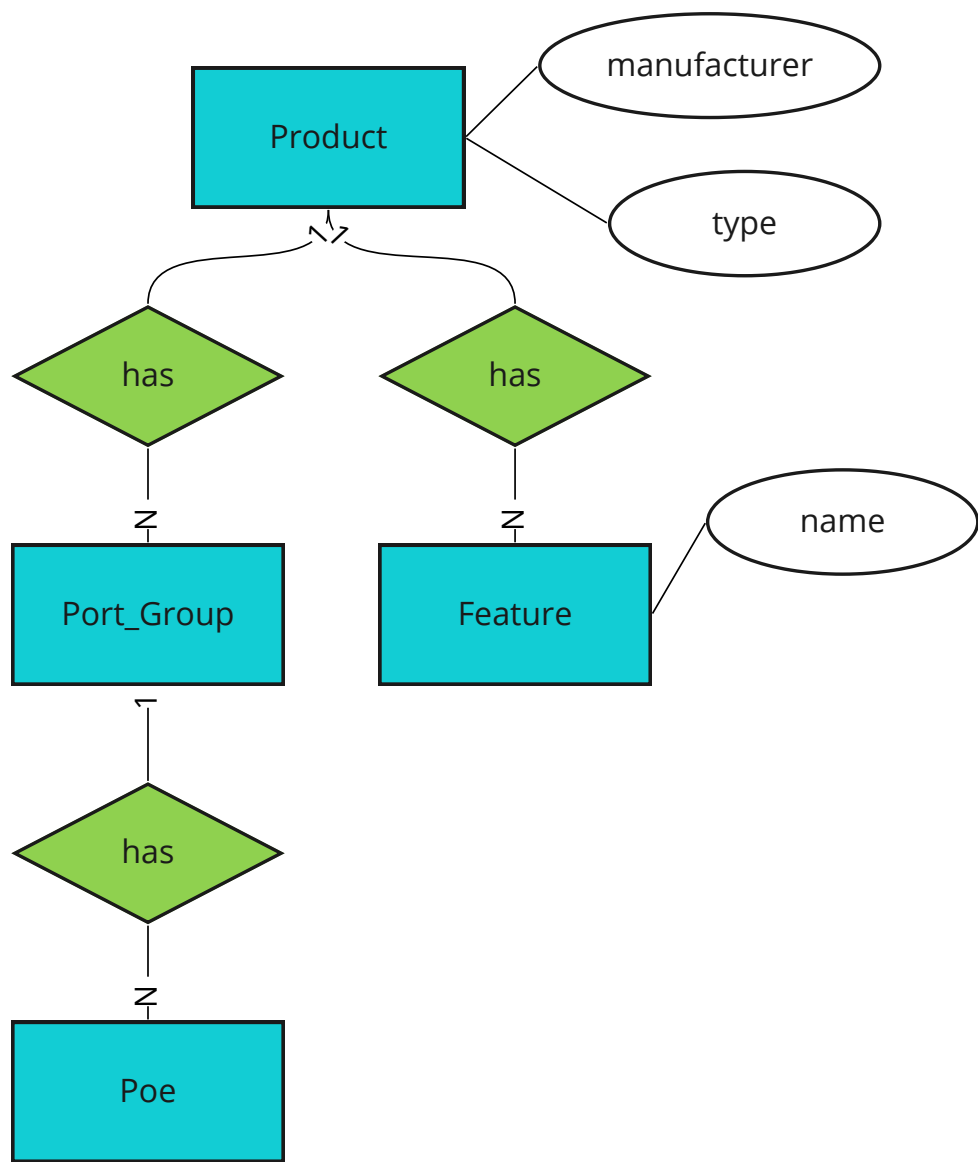

Features can contain features, that's new, but actually also just another attribute in features.


Data model for specific tables

Data model for eav approach

EAV







# data structure

```
Welcome Page | Docker_loacal_ANDREA | edbsHS2023_ANDREA
8.47700024 seconds

Worksheet | Query Builder

ALTER TABLE feature ADD amount_of_rules NUMBER(10);
ALTER TABLE feature ADD amount_of_connections NUMBER(10);
ALTER TABLE feature ADD amount_of_nat_rules NUMBER(10);
ALTER TABLE feature ADD ssl_inspection VARCHAR2(10);
ALTER TABLE feature ADD amount_of_tunnels NUMBER(10);
ALTER TABLE feature ADD amount_of_users NUMBER(10);
ALTER TABLE feature ADD parent_feat_id NUMBER(10);

ALTER TABLE feature ADD CONSTRAINT feat_feat_fk
FOREIGN KEY (parent_feat_id) REFERENCES feature (feat_id);
```



```
Welcome Page | Docker_loacal_ANDREA | edbsHS2023_ANDREA
8.47700024 seconds

Worksheet | Query Builder

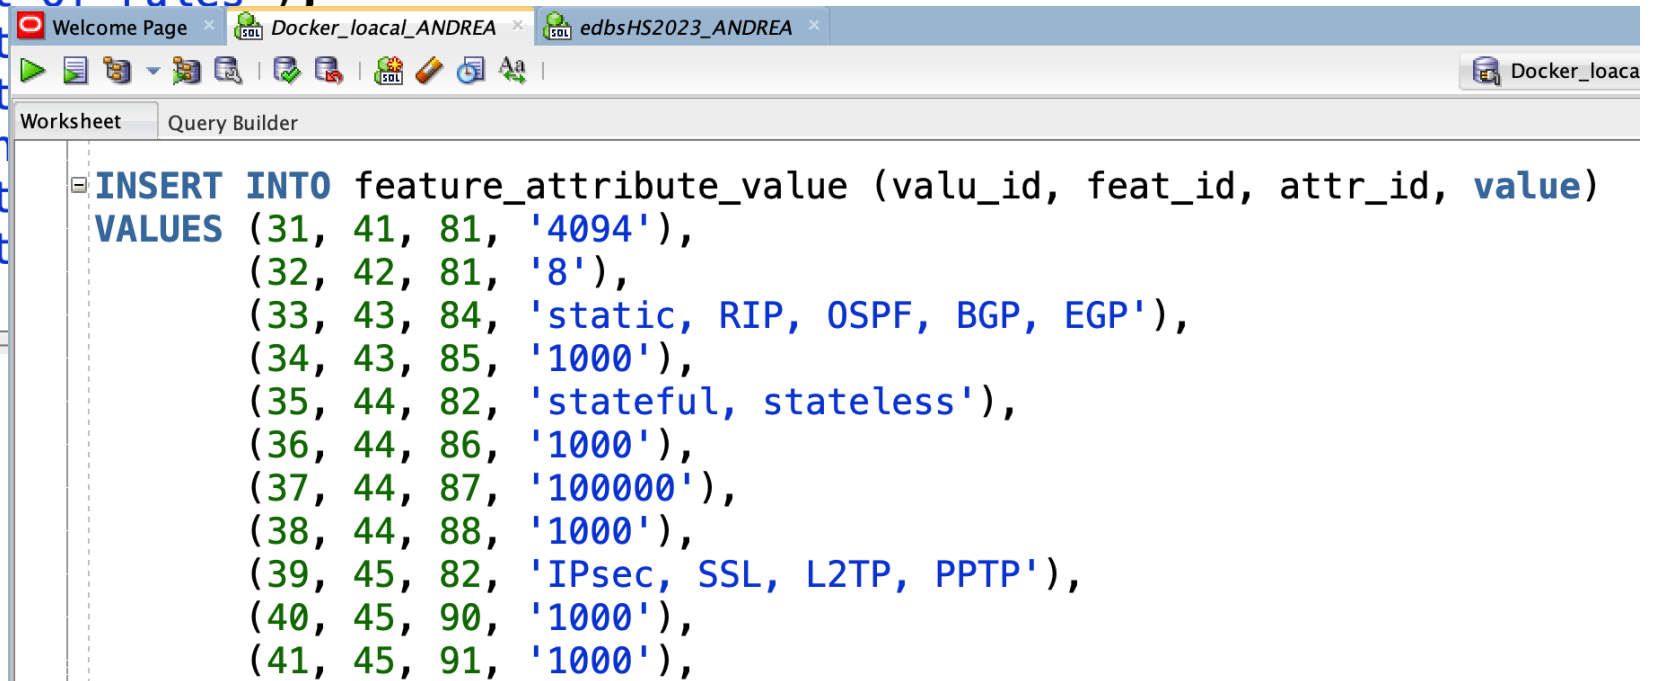
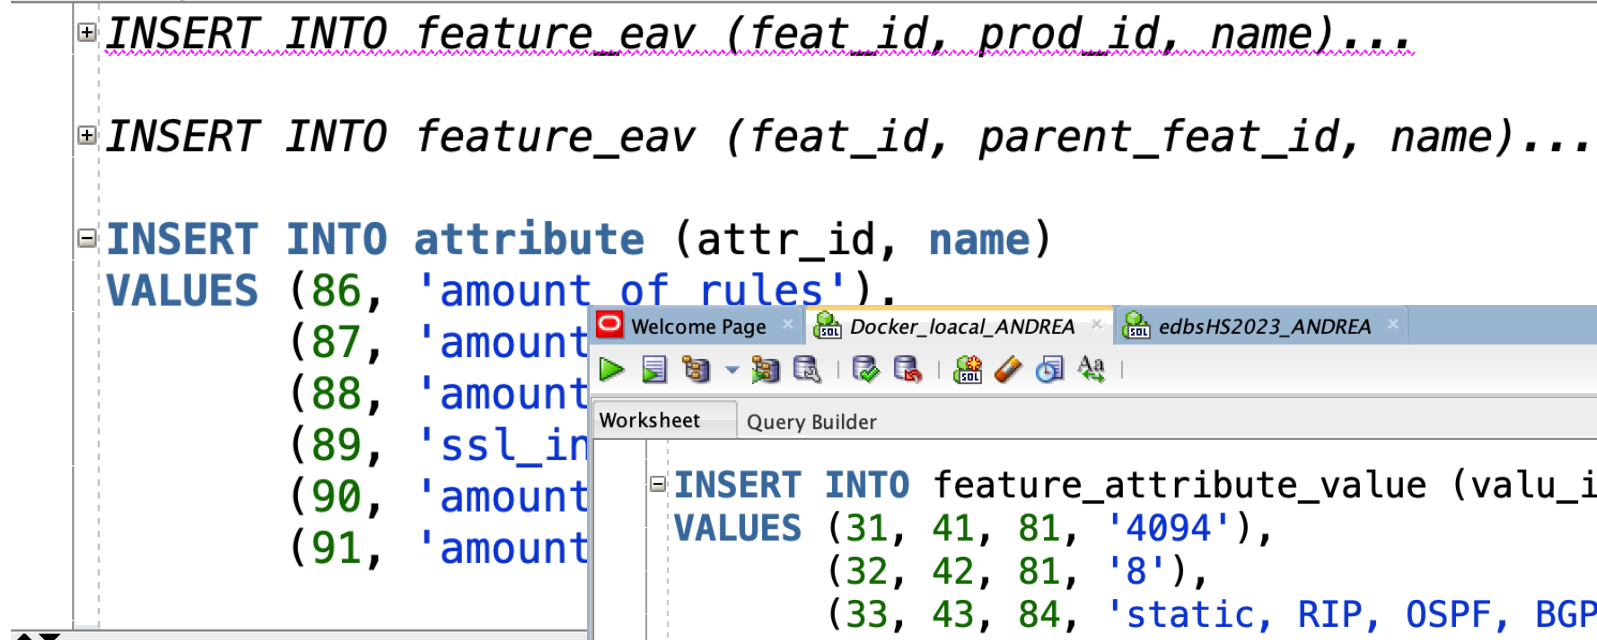
ALTER TABLE feature_eav ADD parent_feat_id NUMBER(10);
ALTER TABLE feature_eav ADD CONSTRAINT feae_feae_fk
FOREIGN KEY (parent_feat_id) REFERENCES feature_eav (feat_id);
```

EAV



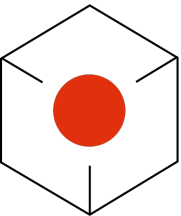


New data

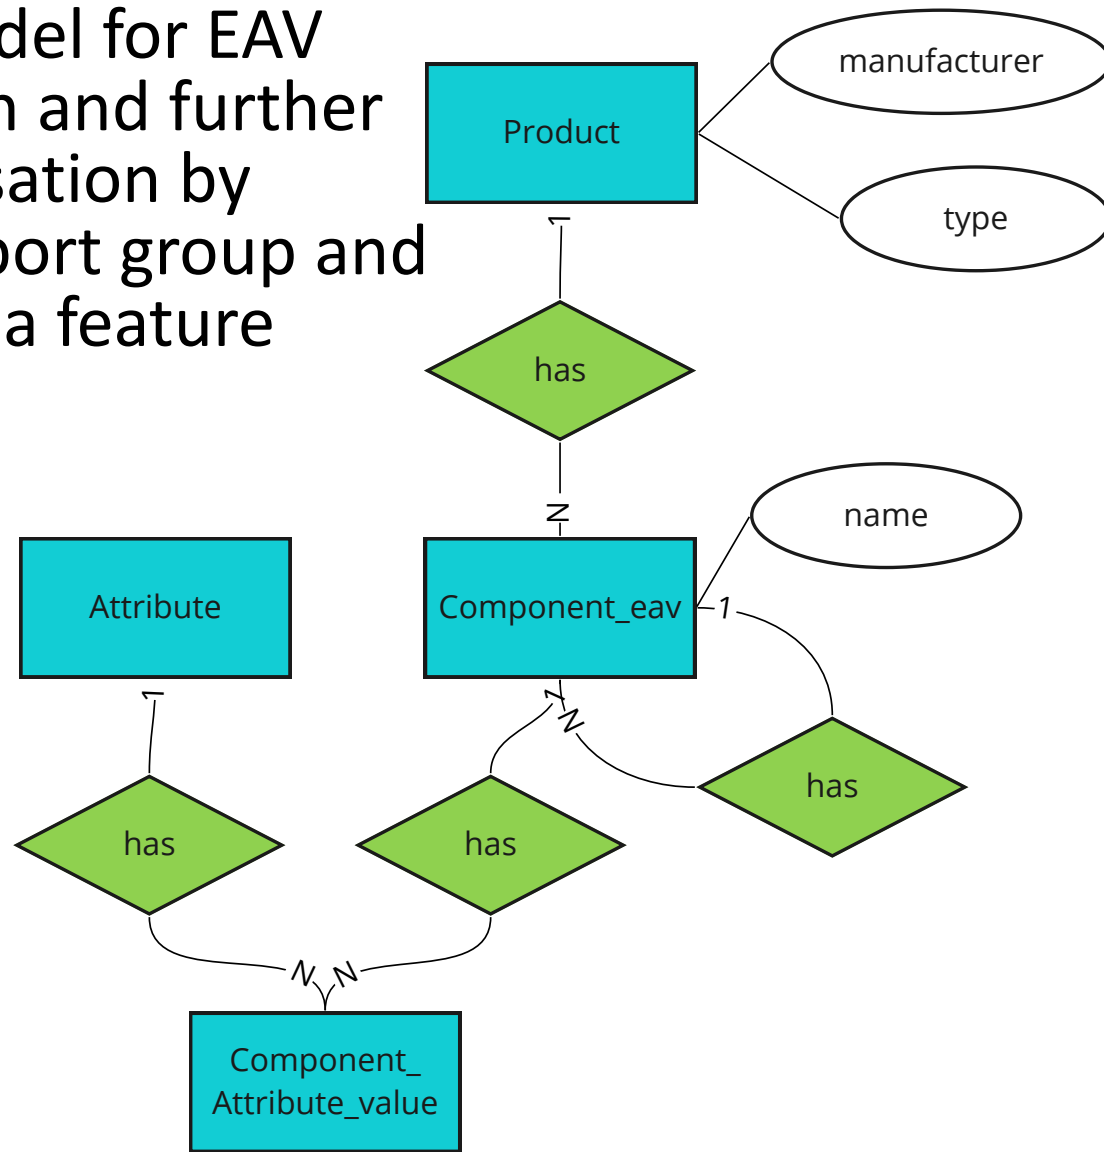
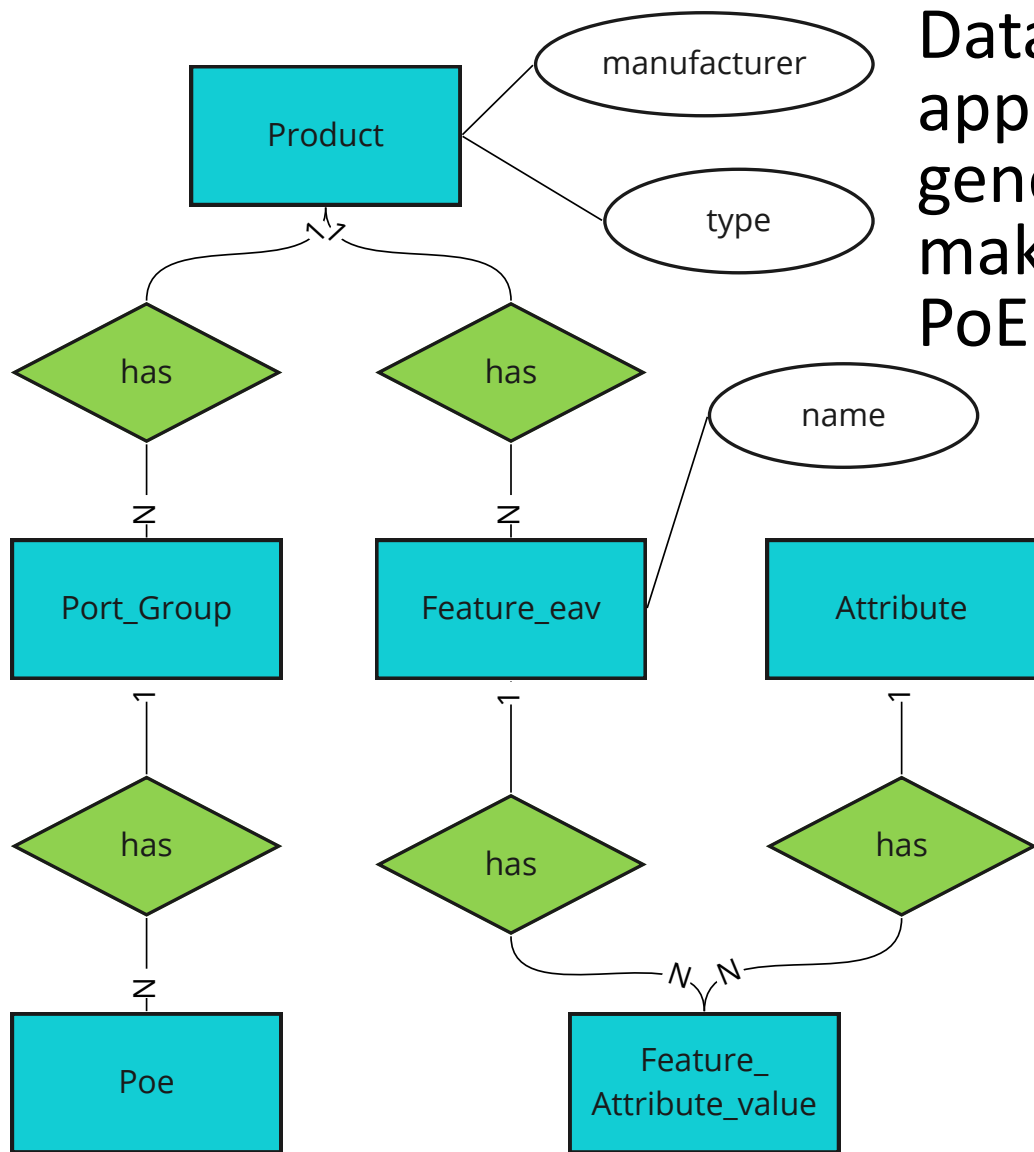


EAV

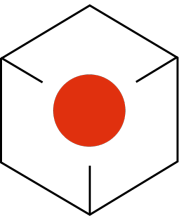




Data model for EAV approach and further generalisation by making port group and PoE also a feature



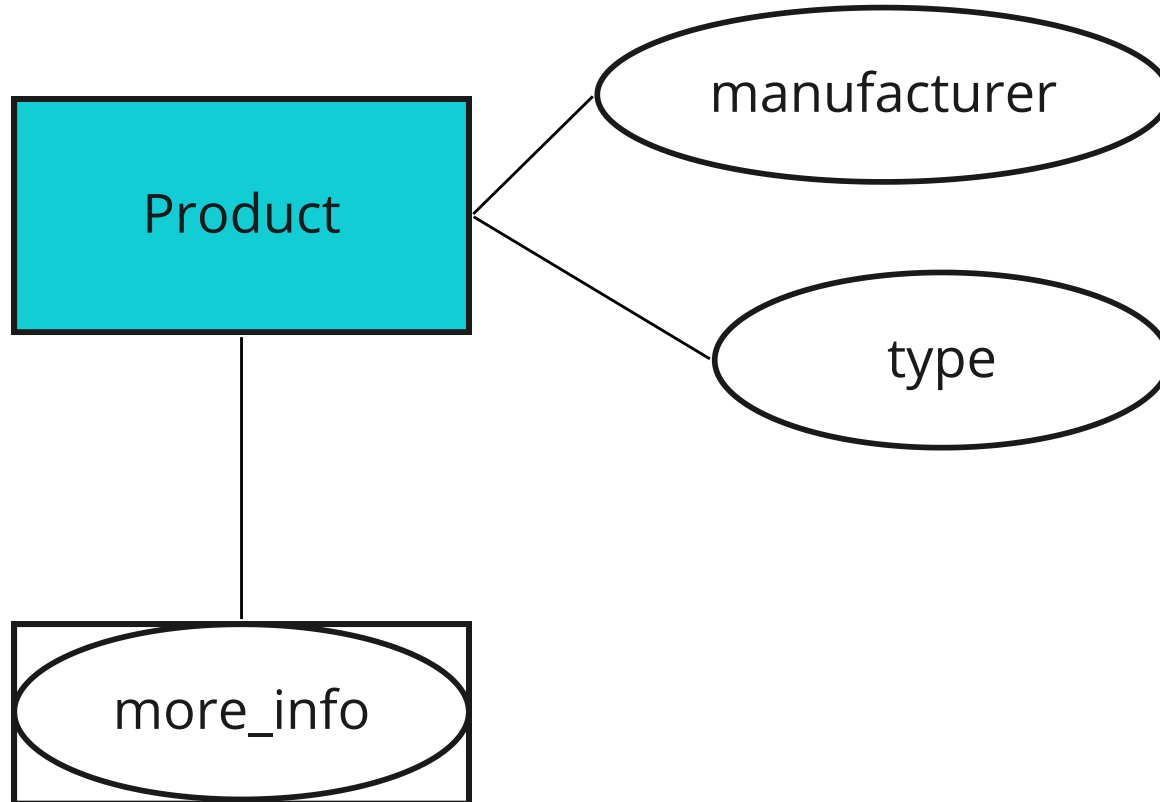
EAV



# Would there be other solutions?

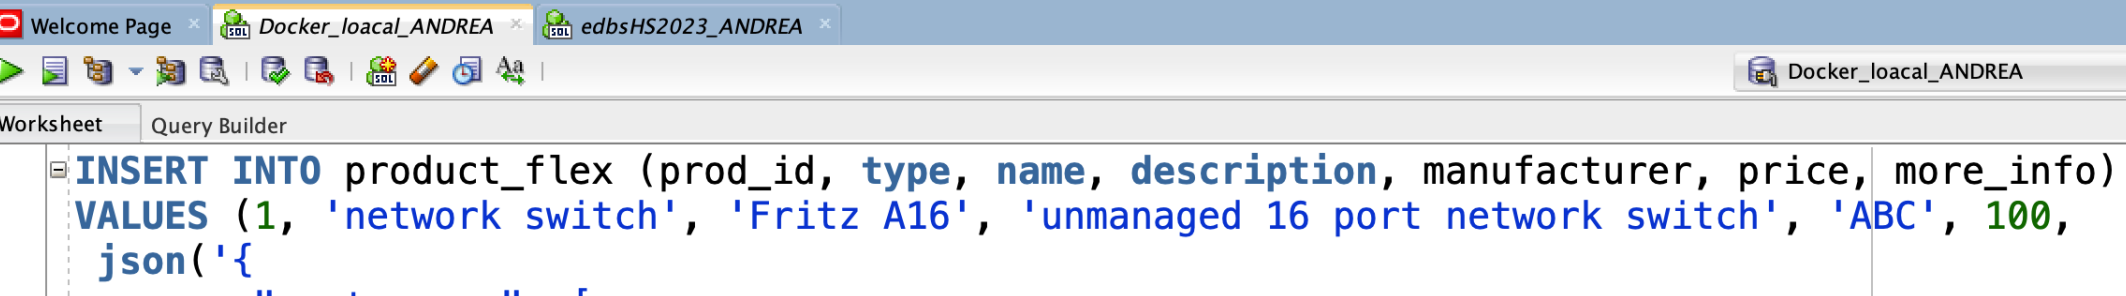
Table		
JSON		

JSON in an attribute  
Explain principle  
Data model  
Demo





# JSON

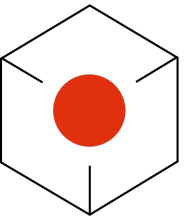


The screenshot shows a SQL IDE window with three tabs: 'Welcome Page', 'Docker\_loacal\_ANDREA', and 'edbsHS2023\_ANDREA'. The 'Worksheet' tab is active, displaying an SQL query. The query is an INSERT statement into a table named 'product\_flex'. The columns specified are 'prod\_id', 'type', 'name', 'description', 'manufacturer', 'price', and 'more\_info'. The values provided are (1, 'network switch', 'Fritz A16', 'unmanaged 16 port network switch', 'ABC', 100, and a JSON object representing port information). The JSON object has a 'port\_group' array containing one object with 'amount': 16, 'type': 'RJ45', and 'speeds': '10/100/1000'.

```

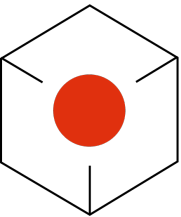
INSERT INTO product_flex (prod_id, type, name, description, manufacturer, price, more_info)
VALUES (1, 'network switch', 'Fritz A16', 'unmanaged 16 port network switch', 'ABC', 100,
json('{
    "port_group": [
        {
            "amount": 16,
            "type": "RJ45",
            "speeds": "10/100/1000"
        }
    ]
}')));

```



# Do we have devices with a port group of type SFP?

```
SELECT pf.prod_id, pf.type product_type, pf.name,  
       jt.*  
FROM product_flex pf,  
     JSON_TABLE(  
       pf.more_info  
       COLUMNS (  
         NESTED port_group[*]  
         COLUMNS (  
           amount NUMBER(3) PATH '$.amount',  
           type VARCHAR2(50) PATH '$.type'  
         )  
       )  
     ) jt  
WHERE jt.type = 'SFP';
```



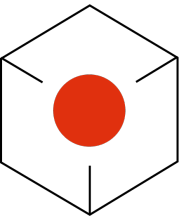
Give me all the details of the network switch “Fritz A16”

```
SELECT p.prod_id, p.type product_type, p.name, p.price
FROM product p INNER JOIN ...


SELECT pc.*
FROM product_collection pc
WHERE pc.json_data.type = 'network switch'
      AND pc.json_data.name = 'Fritz A16';

SELECT pf.*
FROM product_flex pf
WHERE pf.type = 'network switch'
      AND pf.name = 'Fritz A16';
```

Problem  
if JSON has typo



# Summary and vote



- Table with JSON

- EAV

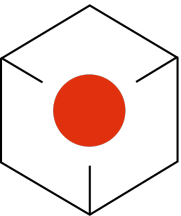
EAV

- Document Database



- Relationale Database





# Dr. Andrea Kennel



Consultant

Lecturer for Databases

Coach for Project Management

University of Applied Sciences  
Northwestern Switzerland

Brugg/Windisch, Switzerland



[andrea.kennel@fhnw.ch](mailto:andrea.kennel@fhnw.ch)

[andrea@infokennel.ch](mailto:andrea@infokennel.ch)

[www.infokennel.ch](http://www.infokennel.ch)